

## E-commerce

```
# import libraries
import pandas as pd
import numpy as np
from scipy.sparse import hstack
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns
from sklearn.model_selection import train_test_split
import re
import nltk
import string
from nltk.tokenize import RegexpTokenizer
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
from bs4 import BeautifulSoup
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import GridSearchCV
import pickle
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, confusion_matrix, \
roc_auc_score, roc_curve, auc, f1_score, \
make_scorer, classification_report
from imblearn.over_sampling import SMOTE
from tensorflow.keras.preprocessing import sequence
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Dense, Input, Dropout, Embedding, \
LSTM, GRU, BatchNormalization
from tensorflow.keras.regularizers import L1L2
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
from tensorflow_addons.metrics import F1Score
import keras.backend as K
import tensorflow as tf
from keras_tuner import RandomSearch, BayesianOptimization
import logging
from textblob import TextBlob

from gensim.models import word2vec
from gensim.models import Word2Vec
```

```

from gensim.models.keyedvectors import KeyedVectors
import gensim
from gensim.models import LdaModel
from gensim import corpora
import pyLDAvis.gensim_models

```

```

import warnings
warnings.filterwarnings('ignore')

```

## Project Task: Week 1

### Class Imbalance Problem:

1. Perform an EDA on the dataset  
*# Get Ecommerce dataset*

```

ecom_df_train = pd.read_csv('train_data.csv')
ecom_df_test = pd.read_csv('test_data.csv')
ecom_df_test_hidden = pd.read_csv('test_data_hidden.csv')

```

*# info*

```

print('train shape ', ecom_df_train.shape)
print('test shape', ecom_df_test.shape)
print('test hidden shape', ecom_df_test_hidden.shape)

```

```

train shape (4000, 8)
test shape (1000, 7)
test hidden shape (1000, 8)

```

```
ecom_df_train.head()
```

		name	brand	\
0	All-New Fire HD 8 Tablet, 8" HD Display, Wi-Fi...		Amazon	
1	Amazon - Echo Plus w/ Built-In Hub - Silver		Amazon	
2	Amazon Echo Show Alexa-enabled Bluetooth Speak...		Amazon	
3	Fire HD 10 Tablet, 10.1 HD Display, Wi-Fi, 16 ...		Amazon	
4	Brand New Amazon Kindle Fire 16gb 7" Ips Displ...		Amazon	

		categories	\
0	Electronics,iPad & Tablets,All Tablets,Fire Ta...		
1	Amazon Echo,Smart Home,Networking,Home & Tools...		
2	Amazon Echo,Virtual Assistant Speakers,Electro...		
3	eBook Readers,Fire Tablets,Electronics Feature...		
4	Computers/Tablets & Networking,Tablets & eBook...		

	primaryCategories	reviews.date	\
0	Electronics	2016-12-26T00:00:00.000Z	
1	Electronics,Hardware	2018-01-17T00:00:00.000Z	
2	Electronics,Hardware	2017-12-20T00:00:00.000Z	
3	Office Supplies,Electronics	2017-08-04T00:00:00.000Z	

4 Electronics 2017-01-23T00:00:00.000Z

```
                                reviews.text \
0  Purchased on Black FridayPros - Great Price (e...
1  I purchased two Amazon in Echo Plus and two do...
2  Just an average Alexa option. Does show a few ...
3  very good product. Exactly what I wanted, and ...
4  This is the 3rd one I've purchased. I've bough...
```

```
                                reviews.title sentiment
0                                Powerful tablet  Positive
1  Amazon Echo Plus AWESOME  Positive
2                                Average  Neutral
3                                Greattttttt  Positive
4                                Very durable!  Positive
```

```
ecom_df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                  4000 non-null  object
1   brand                 4000 non-null  object
2   categories            4000 non-null  object
3   primaryCategories     4000 non-null  object
4   reviews.date          4000 non-null  object
5   reviews.text          4000 non-null  object
6   reviews.title         3990 non-null  object
7   sentiment             4000 non-null  object
dtypes: object(8)
memory usage: 250.1+ KB
```

There are few nan data in reviews.title, others are fine.

a. See what a positive, negative, and neutral review looks like

```
# review sample
```

```
print('\npositive\n-----')
print(ecom_df_train[ecom_df_train['sentiment'] == 'Positive']
['reviews.text'][:1].values)
```

```
print('\nNegative\n-----')
print(ecom_df_train[ecom_df_train['sentiment'] == 'Negative']
['reviews.text'][:1].values)
```

```
print('\nNeutral\n-----')
print(ecom_df_train[ecom_df_train['sentiment'] == 'Neutral']
['reviews.text'][:1].values)
```

positive

```
-----  
['Purchased on Black FridayPros - Great Price (even off sale)Very  
powerful and fast with quad core processors Amazing soundWell  
builtCons -Amazon ads, Amazon need this to subsidize the tablet and  
will remove the adds if you pay them $15.Inability to access other  
apps except the ones from Amazon. There is a way which I was able to  
accomplish to add the Google Play storeNet this is a great tablet for  
the money']
```

Negative

```
-----  
['was cheap, can not run chrome stuff, returned to store.']
```

Neutral

```
-----  
['Just an average Alexa option. Does show a few things on screen but  
still limited.']
```

# NaN check

```
ecom_df_train.isnull().sum()
```

```
name                0  
brand               0  
categories          0  
primaryCategories  0  
reviews.date       0  
reviews.text       0  
reviews.title      10  
sentiment          0  
dtype: int64
```

```
ecom_df_train[ecom_df_train['reviews.title'].isna()][:1]
```

```
name    brand \  
834  Amazon Echo Show Alexa-enabled Bluetooth Speak...  Amazon
```

```
categories \  
primaryCategories \  
834  Computers,Amazon Echo,Virtual Assistant Speake...  
Electronics,Hardware
```

```
reviews.date \  
834  2017-12-29T16:56:05.000Z
```

```
reviews.text reviews.title  
sentiment  
834  Best New Adult Toy in years! Wish I had purcha...  NaN  
Positive
```


As we are mainly focusing on the review texts, so review title is optional. So we can leave it as it is.

b. Check the class count for each class. It's a class imbalance problem.

*# Review type counts*

```
print(ecom_df_train.sentiment.value_counts())  
print('-----')  
print(ecom_df_train.sentiment.value_counts(normalize=True))  
sns.countplot(ecom_df_train.sentiment)  
plt.show()
```

```
Positive      3749  
Neutral       158  
Negative       93  
Name: sentiment, dtype: int64  
-----  
Positive      0.93725  
Neutral       0.03950  
Negative      0.02325  
Name: sentiment, dtype: float64
```



Its highly imbalanced dataset having lots of positive reviews and less negative and neutral reviews.

So while modelling, we need to take this data imbalance into consideration.

**EDA**

`ecom_df_train.columns`

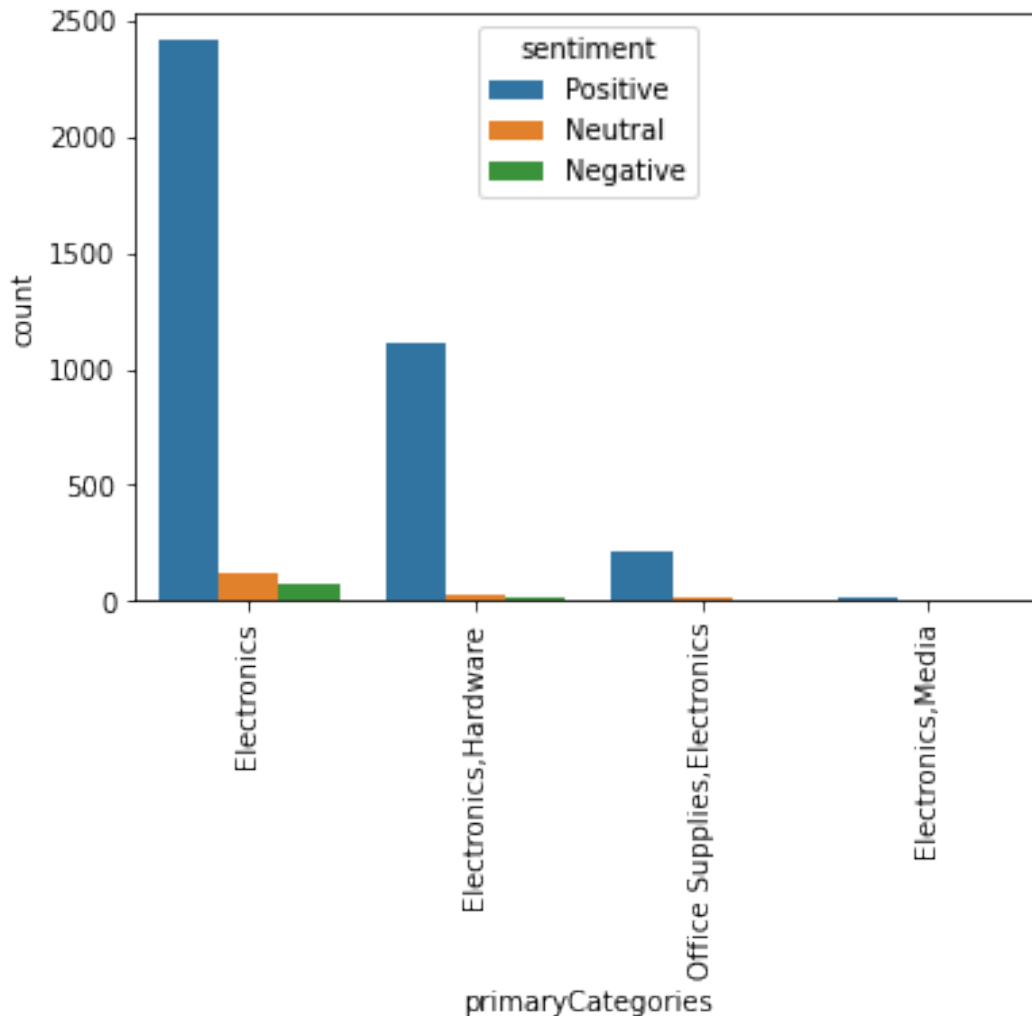
```
Index(['name', 'brand', 'categories', 'primaryCategories',  
      'reviews.date',  
      'reviews.text', 'reviews.title', 'sentiment'],  
      dtype='object')
```

```
# unisug values for each columes
```

```
print('Name unique counts', ecom_df_train.name.nunique())  
print('Brand unique counts', ecom_df_train.brand.nunique())  
print('Categories unique counts', ecom_df_train.categories.nunique())  
print('PrimaryCategories unique counts',  
ecom_df_train.primaryCategories.nunique())  
print('Reviews date unique counts',  
ecom_df_train['reviews.date'].nunique())  
print('Reviews text unique counts',  
ecom_df_train['reviews.text'].nunique())  
print('Reviews title unique counts',  
ecom_df_train['reviews.title'].nunique())  
print('Sentiment unique counts', ecom_df_train['sentiment'].nunique())
```

```
Name unique counts 23  
Brand unique counts 1  
Categories unique counts 23  
PrimaryCategories unique counts 4  
Reviews date unique counts 638  
Reviews text unique counts 3598  
Reviews title unique counts 2606  
Sentiment unique counts 3
```

```
sns.countplot(data=ecom_df_train, x= 'primaryCategories',  
hue='sentiment')  
plt.xticks(rotation=90)  
plt.show()
```



We have most of the reviews from the Electronics category followed by electronic hardware. But everywhere the positive reviews are way high as compared to others.


Here we are mainly focusing on review text to analyse it to get the possible sentiment. Other column values can be ignored.

```
ecom_df_train['reviews.size'] =
ecom_df_train['reviews.text'].apply(len)
```

```
sns.displot(ecom_df_train, x='reviews.size', hue= 'sentiment')
plt.show()
```

WARNING:tensorflow:Detecting that an object or model or tf.train.Checkpoint is being deleted with unrestored values. See the following logs for the specific values in question. To silence these warnings, use `status.expect\_partial()`. See [https://www.tensorflow.org/api\\_docs/python/tf/train/Checkpoint#restore](https://www.tensorflow.org/api_docs/python/tf/train/Checkpoint#restore) for details about the status object returned by the restore function.  
WARNING:tensorflow:Value in checkpoint could not be found in the

```
restored object: (root).optimizer.iter
WARNING:tensorflow:Value in checkpoint could not be found in the
restored object: (root).optimizer.beta_1
WARNING:tensorflow:Value in checkpoint could not be found in the
restored object: (root).optimizer.beta_2
WARNING:tensorflow:Value in checkpoint could not be found in the
restored object: (root).optimizer.decay
WARNING:tensorflow:Value in checkpoint could not be found in the
restored object: (root).optimizer.learning_rate
```



It seems most of the reviews have 100-150 characters, but still there are few reviews with long texts as well.

```
print('Mean review size for positive reviews',
ecom_df_train['reviews.size'][ecom_df_train['sentiment'] ==
'Positive'].median())
print('Mean review size for neutral reviews',
ecom_df_train['reviews.size'][ecom_df_train['sentiment'] ==
'Neutral'].median())
print('Mean review size for negative reviews',
ecom_df_train['reviews.size'][ecom_df_train['sentiment'] ==
'Negative'].median())
```



Mean review size for positive reviews 104.0  
Mean review size for neutral reviews 123.5  
Mean review size for negative reviews 162.0

Avg text size is in between 100 to 150.

### *Feature engineering / text cleaning*

We can now clean the review texts by removing unwanted letters, punctuations and stop words.

It would be helpful on building models that will focus on important words responsible for sentiment prediction.

```
nltk.download('stopwords')
tokenizer = RegexpTokenizer(r'[a-z]+')
stowords_en = stopwords.words("english")
stemmer = SnowballStemmer('english')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.

def preprocess(text):
    text_cln = BeautifulSoup(text, 'lxml').get_text()
    text_cln = re.sub("[^a-zA-Z]", "", text_cln)
    text_cln = ''.join([c for c in text if c not in
string.punctuation])
    text_cln = text_cln.lower()
    text_cln = tokenizer.tokenize(text_cln)
    text_cln = [stemmer.stem(word) for word in text_cln if word not in
stowords_en]
    return ' '.join(text_cln)

ecom_df_train['reviews.clean_text'] =
ecom_df_train['reviews.text'].apply(preprocess)

ecom_df_train['reviews.clean_text'][0]

{"type": "string"}

ecom_df_test['reviews.clean_text'] =
ecom_df_test['reviews.text'].apply(preprocess)

ecom_df_test_hidden['reviews.clean_text'] =
ecom_df_test_hidden['reviews.text'].apply(preprocess)
```

Lets save all the three review texts for future purposes.

We are going to use the preprocessed cleaned review texts and the corresponding sentiments. So lets only take these two columns for each dataframes. For computing sentiment score, we can try with raw and cleaned review texts. So lets keep all the three columns.

```
ecom_df_train = ecom_df_train[['reviews.text', 'reviews.clean_text',
'sentiment']]
ecom_df_test = ecom_df_test[['reviews.text', 'reviews.clean_text']]
ecom_df_test_hidden = ecom_df_test_hidden[['reviews.text',
'reviews.clean_text', 'sentiment']]

ecom_df_train.to_pickle("ecomp_train_processed.pkl")
ecom_df_test_hidden.to_pickle("ecomp_test_hidden_processed.pkl")
ecom_df_test.to_pickle("ecomp_test_processed.pkl")
```

Train test split for modeling

```
X = ecom_df_train['reviews.clean_text']
y = ecom_df_train['sentiment']
```

```
x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, stratify=y, random_state=0)
```

## 2. Convert the reviews in Tf-Idf score.

TFIDF is a function to process the input text to generate n-dim feature vector. As to build model, we need n-dim vectors to represent each data as datapoints in n-dim space, we need to use TF\_IDF vectorization. It uses term frequency and inverse document frequency to give more score to important words used in a review, which is helpful to represent a text review in vector form.

CountVectorizer can also be an option. TFIDF is better than that.

We can customize it as per our need like, we can set params like maximum feature length, min-max word count, n-gram range etc to improve its performance.

Below function can be used to train and use TFIDF vectorizer model, or an already trained model to get the feature vector for a text input.

```
tfidf = TfidfVectorizer(max_features = 5000, max_df=10, min_df=3,
ngram_range=(1, 2))
```

```
def text_transform(x_train, x_test):
    if x_train is not None:
        x_train_vec = tfidf.fit_transform(x_train)
    if x_test is None:
        return x_train_vec
    x_test_vec = tfidf.transform(x_test)
    return x_test_vec if x_train is None else (x_train_vec,
x_test_vec)
```

```
x_train_vec, x_test_vec = text_transform(x_train, x_test)
```

```
print(x_train_vec.shape)
print(x_test_vec.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(3200, 3215)
(800, 3215)
(3200,)
(800,)
```

*5. In case of class imbalance criteria, use the following metrics for evaluating model performance: precision, recall, F1-score, AUC-ROC curve. Use F1-Score as the evaluation criteria for this project.*

To check the model performance, we can use all the classification metrics like accuracy score, precision/recall/f1-score for individual and overall classes as a part of classification report. ROC\_AUC can be also be calculated and AUC curve can be shown for each class label.

It would be helpful to compare the performance between different models as well as for different datasets (train, test etc).

We can mainly focus on the F1-score here as accuracy is not better choice for imbalance data.

```
def getPerformance(y_train, y_pred_train, y_test, y_pred_test,
classes, name):
    print('accuracy score train', accuracy_score(y_train,
y_pred_train))
    print('accuracy score test', accuracy_score(y_test, y_pred_test))

    print('\nTrain classification report: ' + name + '\n')
    print(classification_report(y_train, y_pred_train))
    print('\nTest classification report: ' + name + '\n')
    print(classification_report(y_test, y_pred_test))

    print('\nTrain confusion matrix: ' + name + '\n')
    print(pd.DataFrame(confusion_matrix(y_train, y_pred_train),
columns=classes, index=classes))
    print('\nTest confusion matrix: ' + name + '\n')
    print(pd.DataFrame(confusion_matrix(y_test, y_pred_test),
columns=classes, index=classes))

    train_op = OneVsRest(y_train, y_pred_train)
    test_op = OneVsRest(y_test, y_pred_test)
    for c, y_y_pred in zip(classes, train_op):
        print('\nroc auc score train class {}: {}'.format(c,
roc_auc_score(y_y_pred[0], y_y_pred[1])))

    for c, y_y_pred in zip(classes, train_op):
        fpr, tpr, th = roc_curve(y_y_pred[0], y_y_pred[1])
        auc_val = auc(fpr, tpr)
        plt.title(name + ' Train Receiver Operating Characteristic: '
+ c + ' class')
        plt.plot(fpr, tpr, label='area = {:.3f}'.format(auc_val))
        plt.legend(loc = 'lower right')
```

```

plt.plot([0, 1], [0, 1], 'r--')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

for c, y_y_pred in zip(classes, test_op):
    print('\nroc auc score test class {}: {}'.format(c,
roc_auc_score(y_y_pred[0], y_y_pred[1])))

for c, y_y_pred in zip(classes, test_op):
    fpr, tpr, th = roc_curve(y_y_pred[0], y_y_pred[1])
    auc_val = auc(fpr, tpr)
    plt.title(name + ' Test Receiver Operating Characteristic: ' +
c + ' class')
    plt.plot(fpr, tpr, label='area = {:.3f}'.format(auc_val))
    plt.legend(loc = 'lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()

def OneVsRest(y, y_pred):
    outputs = []
    for c in classes:
        y_one = list(map(lambda x: 1 if x==c else 0, y))
        y_one_red = list(map(lambda x: 1 if x==c else 0, y_pred))
        outputs.append((y_one, y_one_red))
    return outputs

```

*3. Run multinomial Naive Bayes classifier. Everything will be classified as positive because of the class imbalance.*

Multinomial Naive Bayes works as per the naive bayes rule considering the prior probabilities. Its a best choice for textual data.

We can train the model and get sentiment prediction for both train test data and get the model performance using the predicted and actual sentiments.

This steps can be followed for all the models now onwards.

**Multinomial Naive Bayes**

```

mul_nb_model = MultinomialNB()

mul_nb_model.fit(x_train_vec, y_train)
y_pred_train = mul_nb_model.predict(x_train_vec)
y_pred_test = mul_nb_model.predict(x_test_vec)

classes = mul_nb_model.classes_

getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,
'Naive Bayes')

```

accuracy score train 0.9396875  
accuracy score test 0.9375

Train classification report: Naive Bayes

	precision	recall	f1-score	support
Negative	1.00	0.08	0.15	74
Neutral	0.60	0.05	0.09	127
Positive	0.94	1.00	0.97	2999
accuracy			0.94	3200
macro avg	0.85	0.38	0.40	3200
weighted avg	0.93	0.94	0.91	3200

Test classification report: Naive Bayes

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	19
Neutral	0.00	0.00	0.00	31
Positive	0.94	1.00	0.97	750
accuracy			0.94	800
macro avg	0.31	0.33	0.32	800
weighted avg	0.88	0.94	0.91	800

Train confusion matrix: Naive Bayes

	Negative	Neutral	Positive
Negative	6	0	68
Neutral	0	6	121
Positive	0	4	2995

Test confusion matrix: Naive Bayes


	Negative	Neutral	Positive
Negative	0	0	19
Neutral	0	0	31
Positive	0	0	750

roc auc score train class Negative: 0.5405405405405406


roc auc score train class Neutral: 0.5229712174360893

roc auc score train class Positive: 0.529183857305669


Naive Bayes Train Receiver Operating Characteristic: Negative class



Naive Bayes Train Receiver Operating Characteristic: Neutral class



Naive Bayes Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.5


roc auc score test class Neutral: 0.5

roc auc score test class Positive: 0.5


Naive Bayes Test Receiver Operating Characteristic: Negative class



Naive Bayes Test Receiver Operating Characteristic: Neutral class







Now test for the test dataset

```
test_x_vec = text_transform(None,  
ecom_df_test_hidden['reviews.clean_text'])  
test_y = ecom_df_test_hidden['sentiment']  
  
test_y_pred = mul_nb_model.predict(test_x_vec)  
print('Test accuracy score {} \n'.format(accuracy_score(test_y,  
test_y_pred)))  
print('Classification report\n')  
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.939

Classification report

	precision	recall	f1-score	support
Negative	1.00	0.04	0.08	24
Neutral	1.00	0.03	0.05	39
Positive	0.94	1.00	0.97	937
accuracy			0.94	1000
macro avg	0.98	0.36	0.37	1000
weighted avg	0.94	0.94	0.91	1000

```
# save model
pickle.dump(mul_nb_model, open('nb.pkl', 'wb'))
```

We can see all the reviews have been assigned to 'Positive' sentiment. The model is not performing well because of class imbalance.

### Tackling Class Imbalance Problem:

*4. Oversampling or undersampling can be used to tackle the class imbalance problem.*

To handle the class imbalance, we can oversample the data. SMOTE can be used for this by generating synthetic data based on the original data points.

```
sm = SMOTE(random_state=0)
X_blc, y_blc = sm.fit_resample(text_transform(X, None), y)
print('After balancing, train Positive:', sum(y_blc == 'Positive'))
print('After balancing, train Neutral:', sum(y_blc == 'Neutral'))
print('After balancing, train Negative:', sum(y_blc == 'Negative'))
```

```
After balancing, train Positive: 3749
After balancing, train Neutral: 3749
After balancing, train Negative: 3749
```

Now we can see all the three sentiment types have same number of records. Now the model we build, won't face any issue related to class imbalance.

```
x_train_vec, x_test_vec, y_train, y_test = train_test_split(X_blc,
y_blc, test_size=0.2, stratify=y_blc, random_state=0)
print(x_train_vec.shape)
print(x_test_vec.shape)
print(y_train.shape)
print(y_test.shape)

(8997, 4121)
(2250, 4121)
(8997,)
(2250,)
```

### MultinomialNB with data balanced

To check the oversampling impact, we can use naive bayes model and compare its performance against the same model with original data.

```
mul_nb_model_blc = MultinomialNB()
```

Instead of using the model directly, we can do hyper-parameter tuning for the best possible outcome.

```
# Hyper-parameter tuning
```

```
param = {'alpha': [0.001, 0.005, 0.1, 0.5, 1]}
```

```

gs_nb = GridSearchCV(mul_nb_model_blc, param, refit=True,
scoring=make_scorer(f1_score , average='macro'))

gs_nb.fit(x_train_vec, y_train)
y_pred_train = gs_nb.predict(x_train_vec)
y_pred_test = gs_nb.predict(x_test_vec)

classes = gs_nb.classes_

gs_nb.best_score_
0.9333371663199669

gs_nb.best_params_
{'alpha': 0.001}

nb_model = gs_nb.best_estimator_
nb_model
MultinomialNB(alpha=0.001)

getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,
'Naive Bayes Balanced')

accuracy score train 0.9496498832944315
accuracy score test 0.9488888888888889

```

Train classification report: Naive Bayes Balanced

	precision	recall	f1-score	support
Negative	0.89	1.00	0.94	2999
Neutral	0.98	0.94	0.96	2999
Positive	1.00	0.90	0.95	2999
accuracy			0.95	8997
macro avg	0.95	0.95	0.95	8997
weighted avg	0.95	0.95	0.95	8997

Test classification report: Naive Bayes Balanced

	precision	recall	f1-score	support
Negative	0.88	1.00	0.94	750
Neutral	0.98	0.95	0.96	750
Positive	1.00	0.90	0.95	750
accuracy			0.95	2250
macro avg	0.95	0.95	0.95	2250

weighted avg            0.95            0.95            0.95            2250

Train confusion matrix: Naive Bayes Balanced

	Negative	Neutral	Positive
Negative	2999	0	0
Neutral	165	2834	0
Positive	216	72	2711

Test confusion matrix: Naive Bayes Balanced


	Negative	Neutral	Positive
Negative	750	0	0
Neutral	41	709	0
Positive	58	16	676

roc auc score train class Negative: 0.9682394131377126


roc auc score train class Neutral: 0.96648882960987

roc auc score train class Positive: 0.9519839946648883


Naive Bayes Balanced Train Receiver Operating Characteristic: Negative class



Naive Bayes Balanced Train Receiver Operating Characteristic: Neutral class



Naive Bayes Balanced Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.967


roc auc score test class Neutral: 0.9673333333333334

roc auc score test class Positive: 0.9506666666666667


Naive Bayes Balanced Test Receiver Operating Characteristic: Negative class



Naive Bayes Balanced Test Receiver Operating Characteristic: Neutral class



Naive Bayes Balanced Test Receiver Operating Characteristic: Positive class



Lets check model performance on test data.

```
test_x_vec = text_transform(None,
ecom_df_test_hidden['reviews.clean_text'])
test_y = ecom_df_test_hidden['sentiment']

test_y_pred = nb_model.predict(test_x_vec)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.81

Classification report

	precision	recall	f1-score	support
Negative	0.10	0.46	0.16	24
Neutral	0.20	0.33	0.25	39
Positive	0.96	0.84	0.89	937
accuracy			0.81	1000
macro avg	0.42	0.54	0.43	1000
weighted avg	0.91	0.81	0.85	1000

```
# save model
pickle.dump(nb_model, open('nb_blc_bp.pkl', 'wb'))
```

```

ecomp_test_hidden_df = pd.DataFrame.sparse.from_spmatrix(test_x_vec)
ecomp_test_hidden_df = pd.concat([ecomp_test_hidden_df, test_y],
axis=1)
ecomp_test_hidden_df.to_pickle("ecomp_test_hidden_df.pkl")

```

From all the metrics, we can see the model performance has been improved after balanced data is used. Now no longer the score are less for the data with Negative and Neutral setiment.

## 6. Use Tree-based classifiers like Random Forest and XGBoost.

We can go for ensamble models that can enhance the model performance as they use multiple based models with bagging boosting technique.

So we can try out RandomForest and XGBoost which use tree based model (decison tree) as their base model.

We can compare their performance and finally choose the best ML model to go for.

### RandomForest with balanced data

```

rf = RandomForestClassifier(class_weight='balanced')
params = {'n_estimators': [100, 150, 200], 'max_depth': [5, 10, 15],
'min_samples_split': [2, 4, 6]}
gs_rf = GridSearchCV(rf, params, refit=True,
scoring=make_scorer(f1_score , average='macro'))
gs_rf.fit(x_train_vec, y_train)
y_pred_train = gs_rf.predict(x_train_vec)
y_pred_test = gs_rf.predict(x_test_vec)

```

```
classes = gs_rf.classes_
```

```
gs_rf.best_score_
```

```
0.6926930680555026
```

```
gs_rf.best_params_
```

```
{'max_depth': 15, 'min_samples_split': 2, 'n_estimators': 150}
```

```
rf_model = gs_rf.best_estimator_
rf_model
```

```
RandomForestClassifier(class_weight='balanced', max_depth=15,
n_estimators=150)
```

```
getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,
'Random Forest')
```

```
accuracy score train 0.7037901522729799
```

```
accuracy score test 0.6995555555555556
```

Train classification report: Random Forest



	precision	recall	f1-score	support
Negative	1.00	0.62	0.77	2999
Neutral	1.00	0.49	0.66	2999
Positive	0.53	1.00	0.69	2999
accuracy			0.70	8997
macro avg	0.84	0.70	0.71	8997
weighted avg	0.84	0.70	0.71	8997

Test classification report: Random Forest

	precision	recall	f1-score	support
Negative	0.98	0.63	0.77	750
Neutral	0.96	0.50	0.66	750
Positive	0.53	0.97	0.68	750
accuracy			0.70	2250
macro avg	0.82	0.70	0.70	2250
weighted avg	0.82	0.70	0.70	2250

Train confusion matrix: Random Forest

	Negative	Neutral	Positive
Negative	1859	0	1140
Neutral	0	1478	1521
Positive	0	4	2995

Test confusion matrix: Random Forest


	Negative	Neutral	Positive
Negative	475	0	275
Neutral	0	373	377
Positive	10	14	726

roc auc score train class Negative: 0.8099366455485162


roc auc score train class Neutral: 0.7460820273424474

roc auc score train class Positive: 0.777509169723241


Random Forest Train Receiver Operating Characteristic: Negative class



Random Forest Train Receiver Operating Characteristic: Neutral class



Random Forest Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.8133333333333332


roc auc score test class Neutral: 0.7440000000000001

roc auc score test class Positive: 0.7666666666666666


Random Forest Test Receiver Operating Characteristic: Negative class



Random Forest Test Receiver Operating Characteristic: Neutral class



Random Forest Test Receiver Operating Characteristic: Positive class



Lets check for test dataset.

```
test_y_pred = rf_model.predict(test_x_vec)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
```

```
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.921

Classification report

	precision	recall	f1-score	support
Negative	0.21	0.12	0.16	24
Neutral	0.21	0.13	0.16	39
Positive	0.95	0.97	0.96	937
accuracy			0.92	1000
macro avg	0.46	0.41	0.43	1000
weighted avg	0.90	0.92	0.91	1000

*# save model*

```
pickle.dump(rf_model, open('rf_blc_bp.pkl', 'wb'))
```

RandomForest gave better performance than naive bayes for the test data, but for train data, it lacks.

Next we can have XGBoost

*XGBoostClassifier with balanced data*

```
xgc = XGBClassifier(objective='multi:softmax')
params = {
    "gamma": [0.01, 0.05, 0.1],
    "max_depth": [5, 10, 15],
    "n_estimators": [100, 120, 150]
}
gs_xgc = GridSearchCV(xgc, params, refit=True,
    scoring=make_scorer(f1_score, average='macro'))
gs_xgc.fit(x_train_vec, y_train)
y_pred_train = gs_xgc.predict(x_train_vec)
y_pred_test = gs_xgc.predict(x_test_vec)

classes = gs_xgc.classes_

gs_xgc.best_score_

0.920400973797509

gs_xgc.best_params_

{'gamma': 0.1, 'max_depth': 15, 'n_estimators': 150}

xgc_model = gs_xgc.best_estimator_
xgc_model
```

```
XGBClassifier(gamma=0.1, max_depth=15, n_estimators=150,  
              objective='multi:softprob')
```

```
getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,  
'XGBoost')
```

```
accuracy score train 0.942758697343559
```

```
accuracy score test 0.9182222222222223
```

Train classification report: XGBoost

	precision	recall	f1-score	support
Negative	0.99	0.90	0.94	2999
Neutral	0.99	0.95	0.97	2999
Positive	0.87	0.98	0.92	2999
accuracy			0.94	8997
macro avg	0.95	0.94	0.94	8997
weighted avg	0.95	0.94	0.94	8997

Test classification report: XGBoost

	precision	recall	f1-score	support
Negative	0.97	0.89	0.92	750
Neutral	0.95	0.94	0.94	750
Positive	0.85	0.93	0.89	750
accuracy			0.92	2250
macro avg	0.92	0.92	0.92	2250
weighted avg	0.92	0.92	0.92	2250

Train confusion matrix: XGBoost

	Negative	Neutral	Positive
Negative	2698	1	300
Neutral	9	2840	150
Positive	26	29	2944


Test confusion matrix: XGBoost

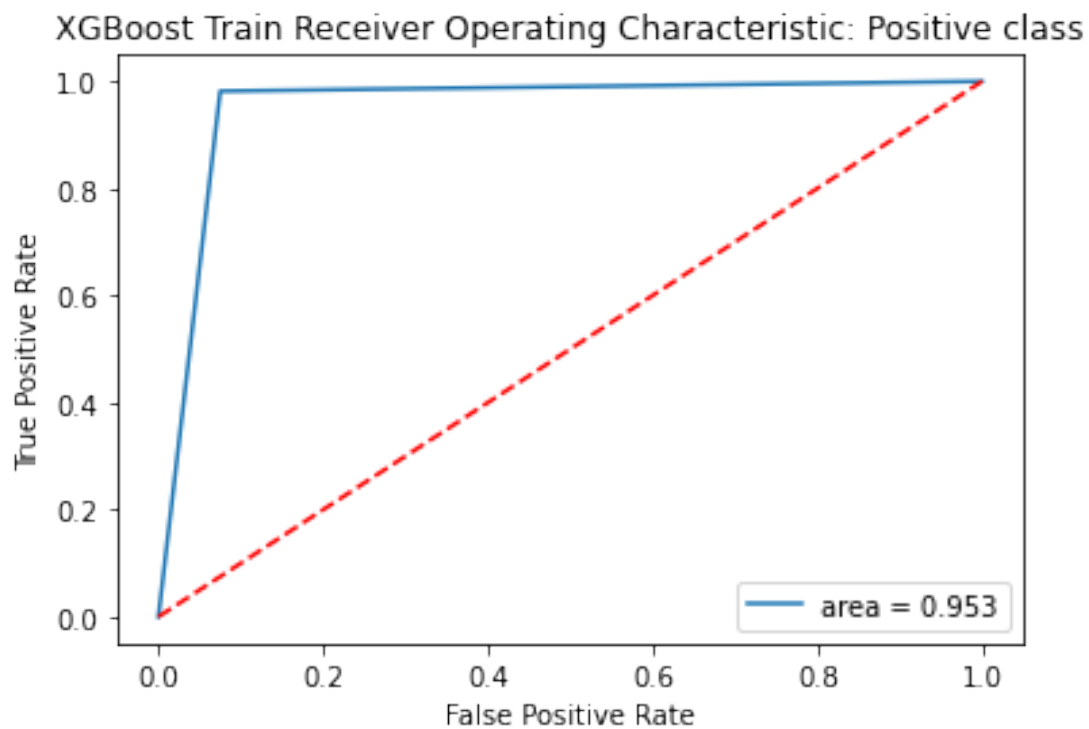

	Negative	Neutral	Positive
Negative	666	2	82
Neutral	6	704	40
Positive	18	36	696

roc auc score train class Negative: 0.9468989663221073

roc auc score train class Neutral: 0.9709903301100367

roc auc score train class Positive: 0.9533177725908636




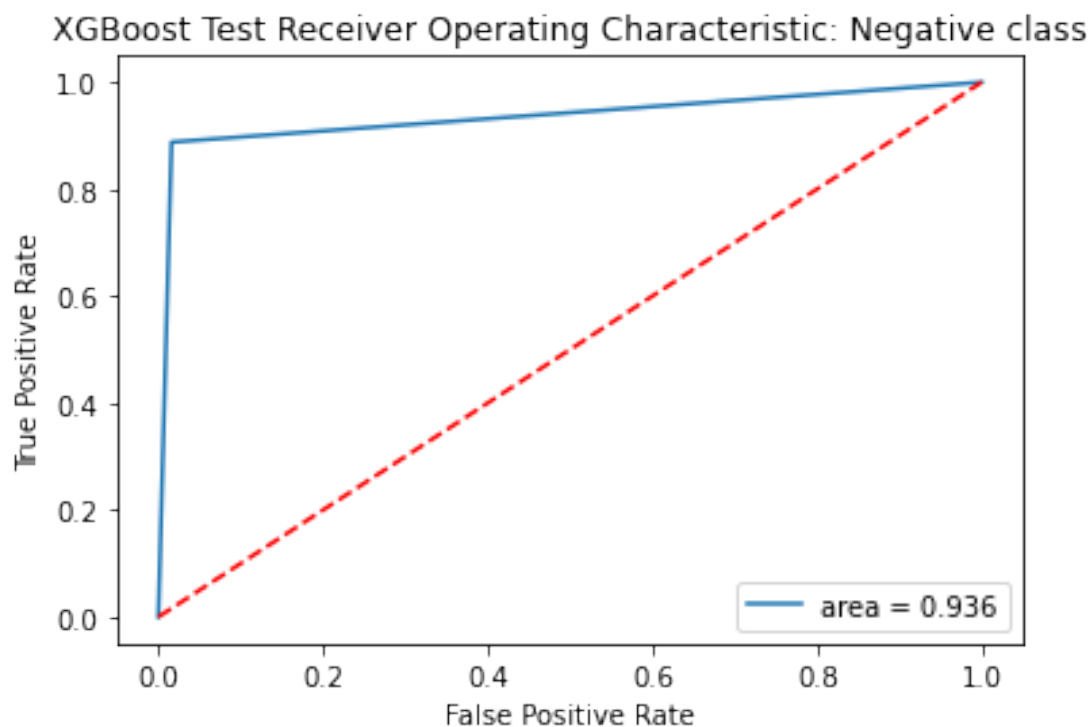



roc auc score test class Negative: 0.9359999999999999

roc auc score test class Neutral: 0.9566666666666668



roc auc score test class Positive: 0.9233333333333333





```
test_y_pred = xgc_model.predict(test_x_vec)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.91

Classification report

	precision	recall	f1-score	support
Negative	0.29	0.33	0.31	24
Neutral	0.27	0.31	0.29	39
Positive	0.96	0.95	0.95	937
accuracy			0.91	1000
macro avg	0.50	0.53	0.52	1000
weighted avg	0.92	0.91	0.91	1000

```
# save model
pickle.dump(xgc_model, open('xgc_blc_bp.pkl', 'wb'))
```

XGBoost performed better than Randomorest and naive bayes for the test data.

Lets save the tfid vectorizer and balanced data for future use.

```
ecomp_train_df = pd.DataFrame.sparse.from_spmatrix(X_blc)
ecomp_train_df = pd.concat([ecomp_train_df, y_blc], axis=1)
ecomp_train_df.to_pickle("ecomp_train_df.pkl")

pickle.dump(tfidf, open('tfidf.pkl', 'wb'))
```

So we can conclude that if it comes to using ML model for the sentiment analysis, we can use XGBoost trained with balanced and tfidf feature transformed textual data.

Below are the macro F1-score for different models with test data.

Multinomial Naive Bayes: 0.37

With balanced data,

Multinomial Naive Bayes: 0.43

RandomForestClassifier: 0.43


XGBClassifier: 0.52

#### Top 10 Important features in XGBoost

We can get the weights representing importance for the features used from XGBoost.

```
feature_names = np.array(tfidf.get_feature_names())
importances = xgc_model.feature_importances_
important_feature = np.argsort(importances)[::-1][:10]
names = feature_names[important_feature]
values = importances[important_feature]

plt.figure(figsize=(10,10))
sns.barplot(names, values)
plt.xticks(rotation=45)
plt.show()
```



## Project Task: Week 2

We can use the tfidf vectorized balanced data for the next ML model, but the neural nets, we can use the actual preprocessed text data for sentiment prediction.

```
ecom_train_df = pd.read_pickle("ecom_train_df.pkl")
print(ecom_train_df.shape)
ecom_train_df.head()
```

```
(11247, 4122)
```

```

      0      1      2      3      4      5      6      7      8      9      ...      4112      4113
4114  \
0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
```

```

0.0
2  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
3  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
4  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0

```

```

      4115  4116  4117  4118  4119  4120  sentiment
0    0.0   0.0   0.0   0.0   0.0   0.0   Positive
1    0.0   0.0   0.0   0.0   0.0   0.0   Positive
2    0.0   0.0   0.0   0.0   0.0   0.0   Neutral
3    0.0   0.0   0.0   0.0   0.0   0.0   Positive
4    0.0   0.0   0.0   0.0   0.0   0.0   Positive

```

[5 rows x 4122 columns]

```

ecomptest_hidden_df = pd.read_pickle("ecomptest_hidden_df.pkl")
print(ecomptest_hidden_df.shape)
ecomptest_hidden_df.head()

```

(1000, 4122)

```

      0      1      2      3      4      5      6      7      8      9  ...  4112  4113
4114  \
0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
2  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
3  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0
4  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
0.0

```

```

      4115  4116  4117  4118  4119  4120  sentiment
0    0.0   0.0   0.0   0.0   0.0   0.0   Positive
1    0.0   0.0   0.0   0.0   0.0   0.0   Positive
2    0.0   0.0   0.0   0.0   0.0   0.0   Positive
3    0.0   0.0   0.0   0.0   0.0   0.0   Positive
4    0.0   0.0   0.0   0.0   0.0   0.0   Positive

```

[5 rows x 4122 columns]

```

test_x_vec = ecomptest_hidden_df.drop(['sentiment'], axis=1)
test_y = ecomptest_hidden_df['sentiment']

```

```

X = ecomptest_train_df.drop(['sentiment'], axis=1)
y = ecomptest_train_df['sentiment']

```

```
x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, stratify=y, random_state=0)
```

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(8997, 4121)
(2250, 4121)
(8997,)
(2250,)
```

### Model Selection:

*2. Use possible ensemble techniques like: XGboost + oversampled\_multinomial\_NB.*

Already covered on week1 (XGBoost and Multinomial NB with balanced data)

*3. Assign a score to the sentence sentiment (engineer a feature called sentiment score). Use this engineered feature in the model and check for improvements. Draw insights on the same.*

Lets use the review texts from the train and test data and compute sentiment score using TextBlob library.

```
ecom_df_train = pd.read_pickle("ecom_train_processed.pkl")
ecom_df_test_hidden =
pd.read_pickle("ecom_test_hidden_processed.pkl")
```

We can create the TFIDF feature vectors with the cleaned text and use it with the sentiment score we get from both plain and clean texts separately.

```
X = ecom_df_train['reviews.clean_text']
X_ = ecom_df_train['reviews.text']
y = ecom_df_train['sentiment']
```

```
test_cln_X = ecom_df_test_hidden['reviews.clean_text']
test_X = ecom_df_test_hidden['reviews.text']
test_y = ecom_df_test_hidden['sentiment']
test_X_vec = text_transform(None, test_cln_X)
```

```
x_cln_train, x_cln_test, x_train, x_test, y_train, y_test =
train_test_split(X, X_, y, test_size=0.2, stratify=y, random_state=0)
```

```
print(x_cln_train.shape)
print(x_cln_test.shape)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(3200,)
(800,)
(3200,)
(800,)
(3200,)
(800,)
```

```
x_train_vec, x_test_vec = text_transform(x_cln_train, x_cln_test)
print(x_train_vec.shape)
print(x_test_vec.shape)
```

```
(3200, 3215)
(800, 3215)
```

Lets try XGBoost model without the sentiment score data, then we can compare with the model with sentiment score.

```
xgb = XGBClassifier(gamma=0.1, max_depth=15, n_estimators=150,
                    objective='multi:softprob')
```

```
xgb.fit(x_train_vec, y_train)
y_pred_train = xgb.predict(x_train_vec)
y_pred_test = xgb.predict(x_test_vec)
```

```
classes = xgb.classes_
```

```
getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,
               'XGBoost')
```

```
accuracy score train 0.9684375
accuracy score test 0.93
```

Train classification report: XGBoost

	precision	recall	f1-score	support
Negative	1.00	0.58	0.74	74
Neutral	0.94	0.47	0.63	127
Positive	0.97	1.00	0.98	2999
accuracy			0.97	3200
macro avg	0.97	0.68	0.78	3200
weighted avg	0.97	0.97	0.96	3200

Test classification report: XGBoost

	precision	recall	f1-score	support
Negative	1.00	0.11	0.19	19
Neutral	0.24	0.13	0.17	31

Positive	0.94	0.98	0.96	750
accuracy			0.93	800
macro avg	0.73	0.41	0.44	800
weighted avg	0.92	0.93	0.91	800

Train confusion matrix: XGBoost

	Negative	Neutral	Positive
Negative	43	1	30
Neutral	0	60	67
Positive	0	3	2996

Test confusion matrix: XGBoost


	Negative	Neutral	Positive
Negative	2	1	16
Neutral	0	4	27
Positive	0	12	738

roc auc score train class Negative: 0.7905405405405406



roc auc score train class Neutral: 0.7355696426329397

roc auc score train class Positive: 0.7582063009394509

XGBoost Train Receiver Operating Characteristic: Negative class










roc auc score test class Negative: 0.5526315789473684

roc auc score test class Neutral: 0.5560635932715298

roc auc score test class Positive: 0.562





```
test_y_pred = xgb.predict(test_X_vec)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.935

Classification report

	precision	recall	f1-score	support
Negative	0.50	0.21	0.29	24
Neutral	0.41	0.18	0.25	39
Positive	0.95	0.99	0.97	937
accuracy			0.94	1000
macro avg	0.62	0.46	0.50	1000
weighted avg	0.92	0.94	0.92	1000

We can set the scores as 0, 5 and 10 for negative, neutral and positive sentiments (score returned from the textBlob polarity can be mapped to these values for better performance)

```
# textBlob
def getScore(text):
    score = round(TextBlob(text).sentiment.polarity, 3)
```

```

if score > 0:
    return 10
elif score < 0:
    return 0
else:
    return 5

# score with raw text
train_score1 = x_train.apply(getScore).values
train_score1 = train_score1.reshape(-1,1)
test_score1 = x_test.apply(getScore).values
test_score1 = test_score1.reshape(-1,1)
# score with cleaned text
train_score2 = x_cln_train.apply(getScore).values
train_score2 = train_score2.reshape(-1,1)
test_score2 = x_cln_test.apply(getScore).values
test_score2 = test_score2.reshape(-1,1)

x_train_vec1 = hstack(blocks= (x_train_vec, train_score1)).tocsr()
x_test_vec1 = hstack(blocks= (x_test_vec, test_score1)).tocsr()

x_train_vec2 = hstack(blocks= (x_train_vec, train_score2)).tocsr()
x_test_vec2 = hstack(blocks= (x_test_vec, test_score2)).tocsr()

test_test_score1 = test_X.apply(getScore).values
test_test_score1 = test_test_score1.reshape(-1,1)
test_test_score2 = test_cln_X.apply(getScore).values
test_test_score2 = test_test_score2.reshape(-1,1)

test_X_vec1 = hstack(blocks= (test_X_vec, test_test_score1)).tocsr()
test_X_vec2 = hstack(blocks= (test_X_vec, test_test_score2)).tocsr()

```

Now lets create an XGBoost model with best parameters with both of these data.

```

xgb1 = XGBClassifier(gamma=0.1, max_depth=15, n_estimators=150,
                    objective='multi:softprob')

xgb1.fit(x_train_vec1, y_train)
y_pred_train1 = xgb1.predict(x_train_vec1)
y_pred_test1 = xgb1.predict(x_test_vec1)

classes = xgb1.classes_

getPerformance(y_train, y_pred_train1, y_test, y_pred_test1, classes,
'XGBoost1')

```

```

accuracy score train 0.9721875
accuracy score test 0.93

```

Train classification report: XGBoost1

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Negative	1.00	0.65	0.79	74
Neutral	0.96	0.53	0.68	127
Positive	0.97	1.00	0.99	2999
accuracy			0.97	3200
macro avg	0.98	0.73	0.82	3200
weighted avg	0.97	0.97	0.97	3200

Test classification report: XGBoost1

	precision	recall	f1-score	support
Negative	1.00	0.11	0.19	19
Neutral	0.26	0.16	0.20	31
Positive	0.95	0.98	0.96	750
accuracy			0.93	800
macro avg	0.74	0.42	0.45	800
weighted avg	0.92	0.93	0.92	800

Train confusion matrix: XGBoost1

	Negative	Neutral	Positive
Negative	48	0	26
Neutral	0	67	60
Positive	0	3	2996

Test confusion matrix: XGBoost1


	Negative	Neutral	Positive
Negative	2	1	16
Neutral	0	5	26
Positive	0	13	737

roc auc score train class Negative: 0.8243243243243243


roc auc score train class Neutral: 0.7632914052030512


roc auc score train class Positive: 0.7855694850190528

XGBoost1 Train Receiver Operating Characteristic: Negative class



XGBoost1 Train Receiver Operating Characteristic: Neutral class






roc auc score test class Negative: 0.5526315789473684


roc auc score test class Neutral: 0.5715424304710769

roc auc score test class Positive: 0.5713333333333334


XGBoost1 Test Receiver Operating Characteristic: Negative class



XGBoost1 Test Receiver Operating Characteristic: Neutral class







```
test_y_pred1 = xgb1.predict(test_X_vec1)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred1)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred1))
```

Test accuracy score 0.937

Classification report

	precision	recall	f1-score	support
Negative	0.71	0.21	0.32	24
Neutral	0.32	0.15	0.21	39
Positive	0.95	0.99	0.97	937
accuracy			0.94	1000
macro avg	0.66	0.45	0.50	1000
weighted avg	0.92	0.94	0.92	1000

```
xgb2 = XGBClassifier(gamma=0.1, max_depth=15, n_estimators=150,
objective='multi:softprob')
```

```
xgb2.fit(x_train_vec2, y_train)
y_pred_train2 = xgb2.predict(x_train_vec2)
y_pred_test2 = xgb2.predict(x_test_vec2)
```

```
classes = xgb2.classes_  
getPerformance(y_train, y_pred_train2, y_test, y_pred_test2, classes,  
'XGBoost2')
```

```
accuracy score train 0.9684375  
accuracy score test 0.93
```

Train classification report: XGBoost2

	precision	recall	f1-score	support
Negative	1.00	0.59	0.75	74
Neutral	0.97	0.46	0.62	127
Positive	0.97	1.00	0.98	2999
accuracy			0.97	3200
macro avg	0.98	0.68	0.78	3200
weighted avg	0.97	0.97	0.96	3200

Test classification report: XGBoost2

	precision	recall	f1-score	support
Negative	1.00	0.11	0.19	19
Neutral	0.20	0.10	0.13	31
Positive	0.94	0.99	0.96	750
accuracy			0.93	800
macro avg	0.71	0.40	0.43	800
weighted avg	0.92	0.93	0.91	800

Train confusion matrix: XGBoost2

	Negative	Neutral	Positive
Negative	44	0	30
Neutral	0	58	69
Positive	0	2	2997



Test confusion matrix: XGBoost2


	Negative	Neutral	Positive
Negative	2	1	16
Neutral	0	3	28
Positive	0	11	739

```
roc auc score train class Negative: 0.7972972972972974
```

roc auc score train class Neutral: 0.7280210417889108

roc auc score train class Positive: 0.7533978988020883






roc auc score test class Negative: 0.5526315789473684


roc auc score test class Neutral: 0.5405847560719829


roc auc score test class Positive: 0.5526666666666666

XGBoost2 Test Receiver Operating Characteristic: Negative class



XGBoost2 Test Receiver Operating Characteristic: Neutral class





```
test_y_pred2 = xgb2.predict(test_X_vec2)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred2)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred2))
```

Test accuracy score 0.937

Classification report

	precision	recall	f1-score	support
Negative	0.50	0.21	0.29	24
Neutral	0.27	0.08	0.12	39
Positive	0.95	0.99	0.97	937
accuracy			0.94	1000
macro avg	0.57	0.43	0.46	1000
weighted avg	0.91	0.94	0.92	1000

```
pickle.dump(xgb1, open('xgc_with_sentiment_score.pkl', 'wb'))
```


From the train and test data performance for all the three XGBClassifier models, we can see that model trained with the sentiment score generated by raw/original review texts has performed better than rest of the models.

Lets find out the feature importances. Hopefully sentiment score should get high importance.

Top 10 important features for XGBoost model with sentiment score

```
feature_names = np.array(tfidf.get_feature_names())  
feature_names = np.append(feature_names, 'score')  
importances = xgb1.feature_importances_  
important_feature = np.argsort(importances)[::-1][:10]  
names = feature_names[important_feature]  
values = importances[important_feature]
```

```
plt.figure(figsize=(10,10))  
sns.barplot(names, values)  
plt.xticks(rotation=45)  
plt.show()
```



We can see 'Score' got the 2nd highest importance among all other features for this XGBoost model. So its concluded that sentiment score can be helpful in improving the model performance.

### 1. Apply multi-class SVM's and neural nets.

We can check performance of multi-class SVM and Neural nets.

#### SVM

```
svm = SGDClassifier(loss='hinge', class_weight='balanced')
svm.fit(x_train, y_train)
y_pred_train = svm.predict(x_train)
y_pred_test = svm.predict(x_test)

classes = svm.classes_

getPerformance(y_train, y_pred_train, y_test, y_pred_test, classes,
'SVM')
```

```
accuracy score train 0.9333111037012337
accuracy score test 0.8911111111111111
```

Train classification report: SVM

	precision	recall	f1-score	support
Negative	0.98	0.89	0.94	2999
Neutral	0.84	0.99	0.91	2999
Positive	1.00	0.91	0.95	2999
accuracy			0.93	8997
macro avg	0.94	0.93	0.93	8997
weighted avg	0.94	0.93	0.93	8997

Test classification report: SVM

	precision	recall	f1-score	support
Negative	0.94	0.88	0.91	750
Neutral	0.79	0.99	0.88	750
Positive	1.00	0.80	0.89	750
accuracy			0.89	2250
macro avg	0.91	0.89	0.89	2250
weighted avg	0.91	0.89	0.89	2250

Train confusion matrix: SVM



	Negative	Neutral	Positive
Negative	2678	321	0
Neutral	15	2984	0
Positive	29	235	2735


Test confusion matrix: SVM



	Negative	Neutral	Positive
Negative	660	90	0
Neutral	8	742	0
Positive	34	113	603

roc auc score train class Negative: 0.9428142714238079

roc auc score train class Neutral: 0.9511503834611537

roc auc score train class Positive: 0.9559853284428143








roc auc score test class Negative: 0.9259999999999999

roc auc score test class Neutral: 0.927

roc auc score test class Positive: 0.902





```
test_y_pred = svm.predict(test_x_vec)
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

Test accuracy score 0.78

Classification report

	precision	recall	f1-score	support
Negative	0.21	0.38	0.27	24
Neutral	0.09	0.41	0.15	39
Positive	0.97	0.81	0.88	937
accuracy			0.78	1000
macro avg	0.42	0.53	0.43	1000
weighted avg	0.91	0.78	0.84	1000

```
# save model
pickle.dump(svm, open('svm_blc_bp.pkl', 'wb'))
```

SVM model performed well, but still XGBoost is the best model till now.

## Neural Nets

As we have text data, we can only use the RNN models for neural nets. We can take clean preprocessed data as input and use LSTM or GRU as RNN models. Additional Dense layers and final Dense layer with Softmax activation will be used.

Also, to feed the actual text data into the LSTM/GRU, we need to encode it by using embedding.

Embedding will convert the text into word tokens and give feature vector for each of them. Keras tokenizer and sequence are also used to support word embedding.

To improve the model, categorical cross-entropy can be used as model loss metrics and f1-score will be the business metric.

```
ecomp_train = pd.read_pickle("ecomp_train_processed.pkl")

ecomp_train['sentiment_score'] =
ecomp_train['sentiment'].map({'Positive': 2, 'Neutral': 1, 'Negative':
0})

ecomp_train['sentiment_score'].value_counts()

2      3749
1       158
0        93
Name: sentiment_score, dtype: int64

ecomp_train.to_pickle("ecomp_train_processed.pkl")

X = ecomp_train['reviews.clean_text']
y = ecomp_train['sentiment_score']
y_label = ecomp_train['sentiment']

x_train, x_test, y_train, y_test, y_train_label, y_test_label =
train_test_split(X, y, y_label,

stratify=y, test_size=0.2,

random_state=0)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_train.shape)

(3200,)
(800,)
(3200,)
(3200,)

ecomp_test_hidden = pd.read_pickle("ecomp_test_hidden_processed.pkl")
ecomp_test_hidden['sentiment_score'] =
```

```

ecomp_test_hidden['sentiment'].map({'Positive': 2, 'Neutral': 1,
'Negative': 0})
ecomp_test_hidden.to_pickle("ecomp_test_hidden_processed.pkl")

test_X = ecomp_test_hidden['reviews.clean_text']
test_y = ecomp_test_hidden['sentiment']

```

### Applying LSTM (LSTM & GRU)

To build the LSTM/GRU model, we will be using embedding layer to convert the text input to feature vector suitable for LSTM/GRU input. After LSTM/GRU layer, we can add series of Dense layers of different filter size. BatchNormalization and Dropout can be used in between for better performance. In the last, one Dense layer with softmax activation will be used for classification.

```

# initiate params
corpus_count = 20000
maxlen = 200
cat_classes = 3
epoch = 10
batch_size = 30

# Prepare input
tokenizer = Tokenizer(num_words=corpus_count)
tokenizer.fit_on_texts(x_train)

pickle.dump(tokenizer, open('keras_tokenizer.pkl', 'wb'))

x_train_seq = tokenizer.texts_to_sequences(x_train)
x_test_seq = tokenizer.texts_to_sequences(x_test)
test_X_seq = tokenizer.texts_to_sequences(test_X)

x_train_seq = sequence.pad_sequences(x_train_seq, maxlen= maxlen)
x_test_seq = sequence.pad_sequences(x_test_seq, maxlen= maxlen)
test_X_seq = sequence.pad_sequences(test_X_seq, maxlen= maxlen)

# Encode Y values
y_train_en = to_categorical(y_train)
y_test_en = to_categorical(y_test)

print(x_train_seq.shape)
print(x_test_seq.shape)
print(y_train.shape)
print(y_test.shape)

(3200, 200)
(800, 200)
(3200,)
(800,)

```

4. Use LSTM for the previous problem (use parameters of LSTM like top-word, embedding-length, Dropout, epochs, number of layers, etc.)

Another variation of LSTM, GRU (Gated Recurrent Units) can be tried as well.

5. Compare the accuracy of neural nets with traditional ML based algorithms.

LSTM

# Model

```
model_lstm = Sequential()
model_lstm.add(Embedding(corpus_count, 150, input_length= maxlen))
model_lstm.add(LSTM(128))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(128, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_lstm.add(Dense(32, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(10, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_lstm.add(Dense(cat_classes, activation='softmax',
kernel_regularizer= L1L2(l1=0.01, l2=0.01)))
```

```
model_lstm.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 200, 150)	3000000
lstm_2 (LSTM)	(None, 128)	142848
dropout_2 (Dropout)	(None, 128)	0
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 32)	4128
dropout_3 (Dropout)	(None, 32)	0
batch_normalization_3 (Batch Normalization)	(None, 32)	128
dense_4 (Dense)	(None, 10)	330

dense\_5 (Dense)

(None, 3)

33

```
=====
Total params: 3,164,491
Trainable params: 3,164,171
Non-trainable params: 320
=====
```

```
model_lstm.compile(loss='categorical_crossentropy',
                    optimizer='adam',
                    metrics=[F1Score(num_classes= 3, average= 'macro')])
```

```
callback = EarlyStopping(monitor='val_loss', patience=3)
```

```
model_lstm.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,
y_test_en),
               batch_size= batch_size, epochs=epoch,
               callbacks=[callback])
```

Epoch 1/10

```
107/107 [=====] - 49s 418ms/step - loss:
13.2222 - f1_score: 0.3224 - val_loss: 6.9888 - val_f1_score: 0.3226
```

Epoch 2/10

```
107/107 [=====] - 44s 408ms/step - loss:
3.3632 - f1_score: 0.3313 - val_loss: 1.5788 - val_f1_score: 0.3226
```

Epoch 3/10

```
107/107 [=====] - 44s 408ms/step - loss:
0.9832 - f1_score: 0.3225 - val_loss: 0.9169 - val_f1_score: 0.3226
```

Epoch 4/10

```
107/107 [=====] - 45s 424ms/step - loss:
0.6539 - f1_score: 0.3225 - val_loss: 0.7162 - val_f1_score: 0.3226
```

Epoch 5/10

```
107/107 [=====] - 44s 416ms/step - loss:
0.5260 - f1_score: 0.3225 - val_loss: 0.5642 - val_f1_score: 0.3226
```

Epoch 6/10

```
107/107 [=====] - 45s 421ms/step - loss:
0.4337 - f1_score: 0.3225 - val_loss: 0.4900 - val_f1_score: 0.3226
```

Epoch 7/10

```
107/107 [=====] - 45s 425ms/step - loss:
0.3891 - f1_score: 0.3225 - val_loss: 0.4437 - val_f1_score: 0.3226
```

Epoch 8/10

```
107/107 [=====] - 45s 420ms/step - loss:
0.3565 - f1_score: 0.3225 - val_loss: 0.4378 - val_f1_score: 0.3226
```

Epoch 9/10

```
107/107 [=====] - 45s 422ms/step - loss:
0.3357 - f1_score: 0.3225 - val_loss: 0.4366 - val_f1_score: 0.3226
```

Epoch 10/10

```
107/107 [=====] - 44s 414ms/step - loss:
0.3179 - f1_score: 0.3225 - val_loss: 0.4108 - val_f1_score: 0.3226
```



<keras.callbacks.History at 0x7fd417153150>

```
def getSentiment(y):  
    d = {2: 'Positive', 1: 'Neutral', 0: 'Negative'}  
    return np.array(list(map(d.get, y.argmax(axis=1))))
```

```
y_pred_train = getSentiment(model_lstm.predict(x_train_seq))  
y_pred_test = getSentiment(model_lstm.predict(x_test_seq))
```

```
100/100 [=====] - 12s 117ms/step  
25/25 [=====] - 3s 116ms/step
```

```
classes = ['Negative', 'Neutral', 'Positive']
```

```
getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,  
classes, 'LSTM')
```

accuracy score train 0.9371875

accuracy score test 0.9375

Train classification report: LSTM

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	74
Neutral	0.00	0.00	0.00	127
Positive	0.94	1.00	0.97	2999
accuracy			0.94	3200
macro avg	0.31	0.33	0.32	3200
weighted avg	0.88	0.94	0.91	3200

Test classification report: LSTM

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	19
Neutral	0.00	0.00	0.00	31
Positive	0.94	1.00	0.97	750
accuracy			0.94	800
macro avg	0.31	0.33	0.32	800
weighted avg	0.88	0.94	0.91	800

Train confusion matrix: LSTM

	Negative	Neutral	Positive
Negative	0	0	74

Neutral	0	0	127
Positive	0	0	2999


Test confusion matrix: LSTM



	Negative	Neutral	Positive
Negative	0	0	19
Neutral	0	0	31
Positive	0	0	750

roc auc score train class Negative: 0.5

roc auc score train class Neutral: 0.5

roc auc score train class Positive: 0.5








roc auc score test class Negative: 0.5

roc auc score test class Neutral: 0.5

roc auc score test class Positive: 0.5





```
test_y_pred = getSentiment(model_lstm.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 4s 115ms/step
Test accuracy score 0.937
```

Classification report

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	24
Neutral	0.00	0.00	0.00	39
Positive	0.94	1.00	0.97	937
accuracy			0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```
model_lstm.save('lstm_model')
```

```
WARNING:absl:Found untraced functions such as
lstm_cell_2_layer_call_fn,
lstm_cell_2_layer_call_and_return_conditional_losses while saving
```

(showing 2 of 2). These functions will not be directly callable after loading.

Neural Nets LSTM didnt perform as expected. Lets try with GRU. Otherwise we can go or simple model.

GRU

# Model

```
model_gru = Sequential()
model_gru.add(Embedding(corpus_count, 150, input_length= maxlen))
model_gru.add(GRU(128))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(128, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_gru.add(Dense(32, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(10, activation='relu', kernel_regularizer=
L1L2(l1=0.01, l2=0.01)))
model_gru.add(Dense(cat_classes, activation='softmax',
kernel_regularizer= L1L2(l1=0.01, l2=0.01)))
```

```
model_gru.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 200, 150)	3000000
gru (GRU)	(None, 128)	107520
dropout_4 (Dropout)	(None, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dense_6 (Dense)	(None, 128)	16512
dense_7 (Dense)	(None, 32)	4128
dropout_5 (Dropout)	(None, 32)	0
batch_normalization_5 (Batch Normalization)	(None, 32)	128
dense_8 (Dense)	(None, 10)	330

dense\_9 (Dense)

(None, 3)

33

```
=====
Total params: 3,129,163
Trainable params: 3,128,843
Non-trainable params: 320
=====
```

```
model_gru.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=[F1Score(num_classes= 3, average= 'macro')])
```

```
model_gru.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,
y_test_en),
              batch_size= batch_size, epochs=epoch,
              callbacks=[callback])
```

```
Epoch 1/10
107/107 [=====] - 51s 441ms/step - loss:
13.4100 - f1_score: 0.3332 - val_loss: 7.6978 - val_f1_score: 0.3226
Epoch 2/10
107/107 [=====] - 48s 448ms/step - loss:
3.9845 - f1_score: 0.3278 - val_loss: 1.9254 - val_f1_score: 0.3226
Epoch 3/10
107/107 [=====] - 38s 359ms/step - loss:
1.1183 - f1_score: 0.3277 - val_loss: 0.9598 - val_f1_score: 0.3226
Epoch 4/10
107/107 [=====] - 51s 479ms/step - loss:
0.6628 - f1_score: 0.3743 - val_loss: 0.7092 - val_f1_score: 0.3226
Epoch 5/10
107/107 [=====] - 50s 473ms/step - loss:
0.5151 - f1_score: 0.4268 - val_loss: 0.5986 - val_f1_score: 0.3226
Epoch 6/10
107/107 [=====] - 36s 333ms/step - loss:
0.4409 - f1_score: 0.3478 - val_loss: 0.5322 - val_f1_score: 0.3226
Epoch 7/10
107/107 [=====] - 52s 483ms/step - loss:
0.3910 - f1_score: 0.3225 - val_loss: 0.4983 - val_f1_score: 0.3226
Epoch 8/10
107/107 [=====] - 52s 489ms/step - loss:
0.3564 - f1_score: 0.3315 - val_loss: 0.4739 - val_f1_score: 0.3226
Epoch 9/10
107/107 [=====] - 43s 402ms/step - loss:
0.3282 - f1_score: 0.3226 - val_loss: 0.4625 - val_f1_score: 0.3226
Epoch 10/10
107/107 [=====] - 41s 379ms/step - loss:
0.3040 - f1_score: 0.3555 - val_loss: 0.4364 - val_f1_score: 0.3561
```

```
<keras.callbacks.History at 0x7fd41c223dd0>
```

```

y_pred_train = getSentiment(model_gru.predict(x_train_seq))
y_pred_test = getSentiment(model_gru.predict(x_test_seq))

100/100 [=====] - 10s 98ms/step
25/25 [=====] - 2s 76ms/step

getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,
classes, 'GRU')

accuracy score train 0.94
accuracy score test 0.93875

```

Train classification report: GRU

	precision	recall	f1-score	support
Negative	0.53	0.11	0.18	74
Neutral	0.67	0.02	0.03	127
Positive	0.94	1.00	0.97	2999
accuracy			0.94	3200
macro avg	0.71	0.37	0.39	3200
weighted avg	0.92	0.94	0.91	3200

Test classification report: GRU

	precision	recall	f1-score	support
Negative	1.00	0.05	0.10	19
Neutral	0.00	0.00	0.00	31
Positive	0.94	1.00	0.97	750
accuracy			0.94	800
macro avg	0.65	0.35	0.36	800
weighted avg	0.90	0.94	0.91	800

Train confusion matrix: GRU

	Negative	Neutral	Positive
Negative	8	1	65
Neutral	6	2	119
Positive	1	0	2998

Test confusion matrix: GRU

	Negative	Neutral	Positive
Negative	1	0	18






Neutral	0	0	31
Positive	0	0	750

roc auc score train class Negative: 0.5529344123394028

roc auc score train class Neutral: 0.5077113082960302

roc auc score train class Positive: 0.5421218349731833








roc auc score test class Negative: 0.5263157894736842

roc auc score test class Neutral: 0.5

roc auc score test class Positive: 0.51





```
test_y_pred = getSentiment(model_gru.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 5s 118ms/step
Test accuracy score 0.938
```

Classification report

	precision	recall	f1-score	support
Negative	1.00	0.04	0.08	24
Neutral	0.00	0.00	0.00	39
Positive	0.94	1.00	0.97	937
accuracy			0.94	1000
macro avg	0.65	0.35	0.35	1000
weighted avg	0.90	0.94	0.91	1000

```
model_gru.save('gru_model')
```

WARNING:absl:Found untraced functions such as gru\_cell\_layer\_call\_fn, gru\_cell\_layer\_call\_and\_return\_conditional\_losses while saving (showing 2 of 2). These functions will not be directly callable after loading.

Both the LSTM and GRU have similar performances.

## Simple LSTM & GRU

As we noticed, for deep learning model, ideally it should perform better than other machine learning models. But with the above architecture, both the LSTM and GRU are not performing as expected. So we can go with simple architecture by removing the kernel regularizers.

### LSTM

#### # Model

```
model_lstm = Sequential()
model_lstm.add(Embedding(corpus_count, 150, input_length= maxlen))
model_lstm.add(LSTM(128))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(128, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dense(32, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(10, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dense(cat_classes, activation='softmax'))

model_lstm.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 200, 150)	3000000
lstm_4 (LSTM)	(None, 128)	142848
dropout_8 (Dropout)	(None, 128)	0
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 128)	16512
dense_17 (Dense)	(None, 32)	4128
dropout_9 (Dropout)	(None, 32)	0
batch_normalization_9 (Batch Normalization)	(None, 32)	128

dense_18 (Dense)	(None, 10)	330
dense_19 (Dense)	(None, 3)	33

```
=====
Total params: 3,164,491
Trainable params: 3,164,171
Non-trainable params: 320
=====
```

```
model_lstm.compile(loss='categorical_crossentropy',
                    optimizer='adam',
                    metrics=[F1Score(num_classes= 3, average= 'macro')])
```

```
callback = EarlyStopping(monitor='val_loss', patience=3)
```

```
model_lstm.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,
y_test_en),
               batch_size= batch_size, epochs=epoch,
               callbacks=[callback])
```

Epoch 1/10

```
107/107 [=====] - 53s 461ms/step - loss:
0.6829 - f1_score: 0.3184 - val_loss: 0.6673 - val_f1_score: 0.3226
```

Epoch 2/10

```
107/107 [=====] - 49s 456ms/step - loss:
0.3009 - f1_score: 0.3408 - val_loss: 0.4563 - val_f1_score: 0.3226
```

Epoch 3/10

```
107/107 [=====] - 51s 477ms/step - loss:
0.2177 - f1_score: 0.4252 - val_loss: 0.3671 - val_f1_score: 0.3226
```

Epoch 4/10

```
107/107 [=====] - 49s 459ms/step - loss:
0.1665 - f1_score: 0.4966 - val_loss: 0.2785 - val_f1_score: 0.3561
```

Epoch 5/10

```
107/107 [=====] - 49s 456ms/step - loss:
0.1278 - f1_score: 0.6643 - val_loss: 0.2508 - val_f1_score: 0.4427
```

Epoch 6/10

```
107/107 [=====] - 51s 473ms/step - loss:
0.0964 - f1_score: 0.7944 - val_loss: 0.3044 - val_f1_score: 0.4510
```

Epoch 7/10

```
107/107 [=====] - 49s 460ms/step - loss:
0.0818 - f1_score: 0.8647 - val_loss: 0.3338 - val_f1_score: 0.5130
```

Epoch 8/10

```
107/107 [=====] - 49s 460ms/step - loss:
0.0522 - f1_score: 0.8957 - val_loss: 0.4166 - val_f1_score: 0.5235
```

```
<keras.callbacks.History at 0x7ff46027f950>
```

```
y_pred_train = getSentiment(model_lstm.predict(x_train_seq))
y_pred_test = getSentiment(model_lstm.predict(x_test_seq))
```

100/100 [=====] - 13s 122ms/step  
25/25 [=====] - 3s 120ms/step

getPerformance(y\_train\_label, y\_pred\_train, y\_test\_label, y\_pred\_test,  
classes, 'LSTM')

accuracy score train 0.9940625  
accuracy score test 0.9375

Train classification report: LSTM

	precision	recall	f1-score	support
Negative	0.99	0.99	0.99	74
Neutral	1.00	0.86	0.92	127
Positive	0.99	1.00	1.00	2999
accuracy			0.99	3200
macro avg	0.99	0.95	0.97	3200
weighted avg	0.99	0.99	0.99	3200

Test classification report: LSTM

	precision	recall	f1-score	support
Negative	0.71	0.26	0.38	19
Neutral	0.33	0.16	0.22	31
Positive	0.95	0.99	0.97	750
accuracy			0.94	800
macro avg	0.67	0.47	0.52	800
weighted avg	0.92	0.94	0.93	800

Train confusion matrix: LSTM

	Negative	Neutral	Positive
Negative	73	0	1
Neutral	1	109	17
Positive	0	0	2999


Test confusion matrix: LSTM

	Negative	Neutral	Positive
Negative	5	2	12
Neutral	0	5	26
Positive	2	8	740



roc auc score train class Negative: 0.9930832944268644

roc auc score train class Neutral: 0.9291338582677166

roc auc score train class Positive: 0.9552238805970149










roc auc score test class Negative: 0.630298537637307

roc auc score test class Neutral: 0.5741432107051471

roc auc score test class Positive: 0.6133333333333333





```
test_y_pred = getSentiment(model_lstm.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 4s 120ms/step
Test accuracy score 0.949
```

Classification report

	precision	recall	f1-score	support
Negative	0.75	0.38	0.50	24
Neutral	0.56	0.23	0.33	39
Positive	0.96	0.99	0.98	937
accuracy			0.95	1000
macro avg	0.76	0.53	0.60	1000
weighted avg	0.94	0.95	0.94	1000

```
model_lstm.save('simple_lstm_model')
```

```
WARNING:absl:Found untraced functions such as
lstm_cell_4_layer_call_fn,
lstm_cell_4_layer_call_and_return_conditional_losses while saving
```

(showing 2 of 2). These functions will not be directly callable after loading.

GRU

# Model

```
model_gru = Sequential()
model_gru.add(Embedding(corpus_count, 150, input_length= maxlen))
model_gru.add(GRU(128))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(128, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dense(32, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(10, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dense(cat_classes, activation='softmax'))
```

```
model_lstm.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 200, 150)	3000000
lstm_4 (LSTM)	(None, 128)	142848
dropout_8 (Dropout)	(None, 128)	0
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 128)	16512
dense_17 (Dense)	(None, 32)	4128
dropout_9 (Dropout)	(None, 32)	0
batch_normalization_9 (Batch Normalization)	(None, 32)	128
dense_18 (Dense)	(None, 10)	330
dense_19 (Dense)	(None, 3)	33
Total params: 3,164,491		

Trainable params: 3,164,171  
Non-trainable params: 320

---

```
model_gru.compile(loss='categorical_crossentropy',  
                  optimizer='adam',  
                  metrics=[F1Score(num_classes= 3, average= 'macro')])
```

```
callback = EarlyStopping(monitor='val_loss', patience=3)
```

```
model_gru.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,  
y_test_en),  
              batch_size= batch_size, epochs=epoch,  
              callbacks=[callback])
```

Epoch 1/10

107/107 [=====] - 47s 412ms/step - loss:  
0.7260 - f1\_score: 0.3000 - val\_loss: 0.5185 - val\_f1\_score: 0.3226

Epoch 2/10

107/107 [=====] - 52s 484ms/step - loss:  
0.3197 - f1\_score: 0.3613 - val\_loss: 0.3577 - val\_f1\_score: 0.3226

Epoch 3/10

107/107 [=====] - 39s 364ms/step - loss:  
0.2198 - f1\_score: 0.4134 - val\_loss: 0.3127 - val\_f1\_score: 0.3226

Epoch 4/10

107/107 [=====] - 39s 362ms/step - loss:  
0.1716 - f1\_score: 0.5501 - val\_loss: 0.3017 - val\_f1\_score: 0.3226

Epoch 5/10

107/107 [=====] - 40s 375ms/step - loss:  
0.1240 - f1\_score: 0.6822 - val\_loss: 0.3053 - val\_f1\_score: 0.3434

Epoch 6/10

107/107 [=====] - 39s 361ms/step - loss:  
0.0956 - f1\_score: 0.7959 - val\_loss: 0.3519 - val\_f1\_score: 0.5189

Epoch 7/10

107/107 [=====] - 39s 362ms/step - loss:  
0.0682 - f1\_score: 0.8647 - val\_loss: 0.3493 - val\_f1\_score: 0.5581

<keras.callbacks.History at 0x7ff460171c10>

```
y_pred_train = getSentiment(model_gru.predict(x_train_seq))  
y_pred_test = getSentiment(model_gru.predict(x_test_seq))
```

100/100 [=====] - 9s 82ms/step  
25/25 [=====] - 2s 87ms/step

```
getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,  
classes, 'GRU')
```

accuracy score train 0.9915625  
accuracy score test 0.93625

Train classification report: GRU

	precision	recall	f1-score	support
Negative	0.96	0.89	0.92	74
Neutral	1.00	0.86	0.92	127
Positive	0.99	1.00	1.00	2999
accuracy			0.99	3200
macro avg	0.98	0.92	0.95	3200
weighted avg	0.99	0.99	0.99	3200

Test classification report: GRU

	precision	recall	f1-score	support
Negative	0.73	0.42	0.53	19
Neutral	0.27	0.13	0.17	31
Positive	0.95	0.98	0.97	750
accuracy			0.94	800
macro avg	0.65	0.51	0.56	800
weighted avg	0.92	0.94	0.93	800

Train confusion matrix: GRU

	Negative	Neutral	Positive
Negative	66	0	8
Neutral	2	109	16
Positive	1	0	2998

Test confusion matrix: GRU


	Negative	Neutral	Positive
Negative	8	1	10
Neutral	0	4	27
Positive	3	10	737

roc auc score train class Negative: 0.9454660994968096


roc auc score train class Neutral: 0.9291338582677166


roc auc score train class Positive: 0.9401317852219397

GRU Train Receiver Operating Characteristic: Negative class



GRU Train Receiver Operating Characteristic: Neutral class






roc auc score test class Negative: 0.7086057011928028

roc auc score test class Neutral: 0.557363983388565


roc auc score test class Positive: 0.6213333333333333




GRU Test Receiver Operating Characteristic: Negative class



GRU Test Receiver Operating Characteristic: Neutral class





```
test_y_pred = getSentiment(model_gru.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 3s 80ms/step
Test accuracy score 0.943
```

Classification report

	precision	recall	f1-score	support
Negative	0.58	0.29	0.39	24
Neutral	0.59	0.26	0.36	39
Positive	0.95	0.99	0.97	937
accuracy			0.94	1000
macro avg	0.71	0.51	0.57	1000
weighted avg	0.93	0.94	0.93	1000

```
model_lstm.save('simple_gru_model')
```

```
WARNING:absl:Found untraced functions such as
lstm_cell_4_layer_call_fn,
lstm_cell_4_layer_call_and_return_conditional_losses while saving
```

(showing 2 of 2). These functions will not be directly callable after loading.

New LSTM model performed better than the GRU and giving the best performance among all the models (based on macro f1-score on the test data).

### *LSTM & GRU with Word2Vec Embedding*

Now instead of training the weights of embedding layer, we can provide our custom embedding weight matrix to the model. We can use word2vec to build the weight matrix.

We can compare the performance of LSTM/GRU with and without w2v embedding.

So first we need to build the w2v model using our existing text inputs and then use the word-vector matrix for the embedding.

#### *# Create corpus*

```
def getCorpus(reviews):  
    sentences = []  
    for review in reviews:  
        sentences.append(review.split())  
    return sentences
```

```
x_train_corpus = getCorpus(x_train)
```

```
x_train_corpus[0]
```

```
['decent', 'camera', 'clariti', 'screen', 'fast', 'download']
```

#### *# params*

```
num_features = 300  
min_wc = 10  
num_workers= -1
```

```
w2v = Word2Vec(sentences= x_train_corpus, size= num_features,  
min_count= min_wc, workers= num_workers)
```

```
WARNING:gensim.models.base_any2vec:EPOCH - 1 : supplied example count  
(0) did not equal expected count (3200)  
WARNING:gensim.models.base_any2vec:EPOCH - 1 : supplied raw word count  
(0) did not equal expected count (47845)  
WARNING:gensim.models.base_any2vec:EPOCH - 2 : supplied example count  
(0) did not equal expected count (3200)  
WARNING:gensim.models.base_any2vec:EPOCH - 2 : supplied raw word count  
(0) did not equal expected count (47845)  
WARNING:gensim.models.base_any2vec:EPOCH - 3 : supplied example count  
(0) did not equal expected count (3200)  
WARNING:gensim.models.base_any2vec:EPOCH - 3 : supplied raw word count  
(0) did not equal expected count (47845)  
WARNING:gensim.models.base_any2vec:EPOCH - 4 : supplied example count  
(0) did not equal expected count (3200)  
WARNING:gensim.models.base_any2vec:EPOCH - 4 : supplied raw word count
```

```

(0) did not equal expected count (47845)
WARNING:gensim.models.base_any2vec:EPOCH - 5 : supplied example count
(0) did not equal expected count (3200)
WARNING:gensim.models.base_any2vec:EPOCH - 5 : supplied raw word count
(0) did not equal expected count (47845)

w2v.save("word2vec.model")

w2v = Word2Vec.load('word2vec.model')

embedding_matrix = w2v.wv.syn0
embedding_matrix.shape

(706, 300)

corpus_count = embedding_matrix.shape[0]
maxlen = 200
cat_classes = 3
epoch = 10
batch_size = 30

tokenizer = Tokenizer(num_words=corpus_count)
tokenizer.fit_on_texts(x_train)

x_train_seq = tokenizer.texts_to_sequences(x_train)
x_test_seq = tokenizer.texts_to_sequences(x_test)
test_X_seq = tokenizer.texts_to_sequences(test_X)

x_train_seq = sequence.pad_sequences(x_train_seq, maxlen= maxlen)
x_test_seq = sequence.pad_sequences(x_test_seq, maxlen= maxlen)
test_X_seq = sequence.pad_sequences(test_X_seq, maxlen= maxlen)

y_train_en = to_categorical(y_train)
y_test_en = to_categorical(y_test)

print(x_train_seq.shape)
print(x_test_seq.shape)
print(y_train.shape)
print(y_test.shape)

(3200, 200)
(800, 200)
(3200,)
(800,)

```

### LSTM with W2V Embedding, dropout and batch normalization

Now Dropout and BatchNormalization can be used to improve the performance.

```

model_lstm = Sequential()
model_lstm.add(Embedding(embedding_matrix.shape[0],
embedding_matrix.shape[1], weights= [embedding_matrix], input_length=
maxlen))

```

```

model_lstm.add(LSTM(128))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(128, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dense(32, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dropout(0.2))
model_lstm.add(BatchNormalization())
model_lstm.add(Dense(10, activation='relu',
kernel_initializer='he_uniform'))
model_lstm.add(Dense(cat_classes, activation='softmax'))

```

```
model_lstm.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
embedding_6 (Embedding)	(None, 200, 300)	211800
lstm_5 (LSTM)	(None, 128)	219648
dropout_12 (Dropout)	(None, 128)	0
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dense_24 (Dense)	(None, 128)	16512
dense_25 (Dense)	(None, 32)	4128
dropout_13 (Dropout)	(None, 32)	0
batch_normalization_13 (Batch Normalization)	(None, 32)	128
dense_26 (Dense)	(None, 10)	330
dense_27 (Dense)	(None, 3)	33

```

Total params: 453,091
Trainable params: 452,771
Non-trainable params: 320

```

```

model_lstm.compile(loss='categorical_crossentropy',
optimizer='adam',
metrics=[F1Score(num_classes= 3, average= 'macro')])

```

```

callback = EarlyStopping(monitor='val_loss', patience=3)

model_lstm.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,
y_test_en),
               batch_size= batch_size, epochs=epoch,
               callbacks=[callback])

Epoch 1/10
107/107 [=====] - 65s 575ms/step - loss:
0.7329 - f1_score: 0.3061 - val_loss: 0.6617 - val_f1_score: 0.3226
Epoch 2/10
107/107 [=====] - 56s 521ms/step - loss:
0.2880 - f1_score: 0.3371 - val_loss: 0.3287 - val_f1_score: 0.3226
Epoch 3/10
107/107 [=====] - 55s 514ms/step - loss:
0.2430 - f1_score: 0.3520 - val_loss: 0.2795 - val_f1_score: 0.3226
Epoch 4/10
107/107 [=====] - 57s 530ms/step - loss:
0.1952 - f1_score: 0.5066 - val_loss: 0.2624 - val_f1_score: 0.3226
Epoch 5/10
107/107 [=====] - 55s 516ms/step - loss:
0.1669 - f1_score: 0.5806 - val_loss: 0.2450 - val_f1_score: 0.3224
Epoch 6/10
107/107 [=====] - 57s 535ms/step - loss:
0.1323 - f1_score: 0.7065 - val_loss: 0.2645 - val_f1_score: 0.3226
Epoch 7/10
107/107 [=====] - 56s 520ms/step - loss:
0.1066 - f1_score: 0.7564 - val_loss: 0.2923 - val_f1_score: 0.4331
Epoch 8/10
107/107 [=====] - 55s 514ms/step - loss:
0.1002 - f1_score: 0.7976 - val_loss: 0.2841 - val_f1_score: 0.5034

<keras.callbacks.History at 0x7ff45a48ec10>

y_pred_train = getSentiment(model_lstm.predict(x_train_seq))
y_pred_test = getSentiment(model_lstm.predict(x_test_seq))

100/100 [=====] - 16s 147ms/step
25/25 [=====] - 4s 146ms/step

getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,
classes, 'LSTM with W2V Embedding')

accuracy score train 0.984375
accuracy score test 0.94

```

Train classification report: LSTM with W2V Embedding

	precision	recall	f1-score	support
Negative	0.96	0.88	0.92	74

Neutral	0.95	0.72	0.82	127
Positive	0.99	1.00	0.99	2999
accuracy			0.98	3200
macro avg	0.96	0.86	0.91	3200
weighted avg	0.98	0.98	0.98	3200

Test classification report: LSTM with W2V Embedding

	precision	recall	f1-score	support
Negative	0.50	0.21	0.30	19
Neutral	0.50	0.16	0.24	31
Positive	0.95	0.99	0.97	750
accuracy			0.94	800
macro avg	0.65	0.45	0.50	800
weighted avg	0.92	0.94	0.93	800

Train confusion matrix: LSTM with W2V Embedding

	Negative	Neutral	Positive
Negative	65	1	8
Neutral	2	91	34
Positive	1	4	2994

Test confusion matrix: LSTM with W2V Embedding


	Negative	Neutral	Positive
Negative	4	1	14
Neutral	1	5	25
Positive	3	4	743

roc auc score train class Negative: 0.9387093427400529


roc auc score train class Neutral: 0.8574541792754267

roc auc score train class Positive: 0.8946887768559669

LSTM with W2V Embedding Train Receiver Operating Characteristic: Negative class




LSTM with W2V Embedding Train Receiver Operating Characteristic: Neutral class





LSTM with W2V Embedding Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.6027023384325089


roc auc score test class Neutral: 0.5773941859977348

roc auc score test class Positive: 0.6053333333333334


LSTM with W2V Embedding Test Receiver Operating Characteristic: Negative class



LSTM with W2V Embedding Test Receiver Operating Characteristic: Neutral class



LSTM with W2V Embedding Test Receiver Operating Characteristic: Positive class



```
test_y_pred = getSentiment(model_lstm.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 5s 146ms/step
Test accuracy score 0.948
```

Classification report

precision	recall	f1-score	support
-----------	--------	----------	---------

Negative	0.69	0.38	0.49	24
Neutral	0.55	0.28	0.37	39
Positive	0.96	0.99	0.97	937
accuracy			0.95	1000
macro avg	0.73	0.55	0.61	1000
weighted avg	0.94	0.95	0.94	1000

```
model_lstm.save('lstm_w2v_model')
```

WARNING:absl:Found untraced functions such as lstm\_cell\_5\_layer\_call\_fn, lstm\_cell\_5\_layer\_call\_and\_return\_conditional\_losses while saving (showing 2 of 2). These functions will not be directly callable after loading.

So we can see, with our provided embedding weight matrix, the lstm model is not performing better than the existing lstm.

#### GRU with W2V Embedding, dropout and batch normalization

We can use Dropout and BatchNormalization here as well.

```
model_gru = Sequential()
model_gru.add(Embedding(embedding_matrix.shape[0],
embedding_matrix.shape[1], weights= [embedding_matrix], input_length=
maxlen))
model_gru.add(GRU(128))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(128, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dense(32, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dropout(0.2))
model_gru.add(BatchNormalization())
model_gru.add(Dense(10, activation='relu',
kernel_initializer='he_uniform'))
model_gru.add(Dense(cat_classes, activation='softmax'))
```

```
model_gru.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
embedding_7 (Embedding)	(None, 200, 300)	211800
gru_1 (GRU)	(None, 128)	165120

dropout_14 (Dropout)	(None, 128)	0
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dense_28 (Dense)	(None, 128)	16512
dense_29 (Dense)	(None, 32)	4128
dropout_15 (Dropout)	(None, 32)	0
batch_normalization_15 (Batch Normalization)	(None, 32)	128
dense_30 (Dense)	(None, 10)	330
dense_31 (Dense)	(None, 3)	33

```

=====
Total params: 398,563
Trainable params: 398,243
Non-trainable params: 320

```

```

model_gru.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=[F1Score(num_classes= 3, average= 'macro')])

```

```

model_gru.fit(x_train_seq, y_train_en, validation_data=(x_test_seq,
y_test_en),
              batch_size= batch_size, epochs=epoch,
              callbacks=[callback])

```

```

Epoch 1/10
107/107 [=====] - 50s 438ms/step - loss:
0.4661 - f1_score: 0.3224 - val_loss: 0.4255 - val_f1_score: 0.3226
Epoch 2/10
107/107 [=====] - 47s 440ms/step - loss:
0.2576 - f1_score: 0.3311 - val_loss: 0.2841 - val_f1_score: 0.3226
Epoch 3/10
107/107 [=====] - 46s 430ms/step - loss:
0.2127 - f1_score: 0.4114 - val_loss: 0.2792 - val_f1_score: 0.3226
Epoch 4/10
107/107 [=====] - 46s 428ms/step - loss:
0.1906 - f1_score: 0.5146 - val_loss: 0.2703 - val_f1_score: 0.3226
Epoch 5/10
107/107 [=====] - 46s 430ms/step - loss:
0.1663 - f1_score: 0.5365 - val_loss: 0.2734 - val_f1_score: 0.3561
Epoch 6/10
107/107 [=====] - 47s 439ms/step - loss:

```

```
0.1466 - f1_score: 0.6178 - val_loss: 0.2906 - val_f1_score: 0.3428
Epoch 7/10
107/107 [=====] - 46s 428ms/step - loss:
0.1247 - f1_score: 0.6844 - val_loss: 0.3106 - val_f1_score: 0.4522
```

```
<keras.callbacks.History at 0x7ff45a5241d0>
```

```
y_pred_train = getSentiment(model_gru.predict(x_train_seq))
y_pred_test = getSentiment(model_gru.predict(x_test_seq))
```

```
100/100 [=====] - 11s 109ms/step
25/25 [=====] - 3s 109ms/step
```

```
getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,
classes, 'GRU with W2V Embedding')
```

```
accuracy score train 0.9759375
accuracy score test 0.92625
```

Train classification report: GRU with W2V Embedding

	precision	recall	f1-score	support
Negative	0.82	0.62	0.71	74
Neutral	0.89	0.67	0.77	127
Positive	0.98	1.00	0.99	2999
accuracy			0.98	3200
macro avg	0.90	0.76	0.82	3200
weighted avg	0.97	0.98	0.97	3200

Test classification report: GRU with W2V Embedding

	precision	recall	f1-score	support
Negative	0.57	0.21	0.31	19
Neutral	0.13	0.06	0.09	31
Positive	0.94	0.98	0.96	750
accuracy			0.93	800
macro avg	0.55	0.42	0.45	800
weighted avg	0.90	0.93	0.91	800

Train confusion matrix: GRU with W2V Embedding

	Negative	Neutral	Positive
Negative	46	7	21
Neutral	6	85	36

Positive            4            3            2992

Test confusion matrix: GRU with W2V Embedding


	Negative	Neutral	Positive
Negative	4	0	15
Neutral	1	2	28
Positive	2	13	735

roc auc score train class Negative: 0.8092113226470232


roc auc score train class Neutral: 0.8330185947713256

roc auc score train class Positive: 0.8570418995386524


GRU with W2V Embedding Train Receiver Operating Characteristic: Negative class



GRU with W2V Embedding Train Receiver Operating Characteristic: Neutral class



GRU with W2V Embedding Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.6033425432980659


roc auc score test class Neutral: 0.5238055287554008

roc auc score test class Positive: 0.56

GRU with W2V Embedding Test Receiver Operating Characteristic: Negative class




GRU with W2V Embedding Test Receiver Operating Characteristic: Neutral class





GRU with W2V Embedding Test Receiver Operating Characteristic: Positive class



```
test_y_pred = getSentiment(model_gru.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 4s 107ms/step
Test accuracy score 0.938
```

Classification report

	precision	recall	f1-score	support
Negative	0.45	0.21	0.29	24
Neutral	0.50	0.23	0.32	39
Positive	0.95	0.99	0.97	937
accuracy			0.94	1000
macro avg	0.64	0.48	0.52	1000
weighted avg	0.92	0.94	0.93	1000

```
model_gru.save('gru_w2v_model')
```

WARNING:absl:Found untraced functions such as gru\_cell\_1\_layer\_call\_fn, gru\_cell\_1\_layer\_call\_and\_return\_conditional\_losses while saving (showing 2 of 2). These functions will not be directly callable after loading.

LSTM with provided embedding matrix is performing much better than the same with GRU and its performance is similar to the simple LSTM.

So LSTM with w2v embedding gave the best performance till now.

So we can let the embedding layer to train itself by going for simple LSTM/GRU models, but hyper parameter tuning can be done on those models.

*6. Find the best setting of LSTM (Neural Net) and GRU that can best classify the reviews as positive, negative, and neutral.*

We can take the input data and use keras tuner RandomSearch to find best model structure and params.

As we got better performance by using the self trained embedding layer without explicit embedding matrix weights, we can use the same input here as well.

```
# initiate params
corpus_count = 20000
maxlen = 200
cat_classes = 3
epoch = 5
batch_size = 30
trials = 3

# Prepare input
tokenizer = Tokenizer(num_words=corpus_count)
tokenizer.fit_on_texts(x_train)

pickle.dump(tokenizer, open('keras_tokenizer.pkl', 'wb'))

x_train_seq = tokenizer.texts_to_sequences(x_train)
x_test_seq = tokenizer.texts_to_sequences(x_test)
test_X_seq = tokenizer.texts_to_sequences(test_X)

x_train_seq = sequence.pad_sequences(x_train_seq, maxlen= maxlen)
x_test_seq = sequence.pad_sequences(x_test_seq, maxlen= maxlen)
test_X_seq = sequence.pad_sequences(test_X_seq, maxlen= maxlen)

# Encode Y values
y_train_en = to_categorical(y_train)
y_test_en = to_categorical(y_test)

print(x_train_seq.shape)
print(x_test_seq.shape)
print(y_train.shape)
print(y_test.shape)

(3200, 200)
(800, 200)
(3200,)
(800,)
```

```

# Model builder
def modelBuilder(hy_param):
    # input
    inputs = Input(shape=(200,))
    # Tokenization/embedding param
    corpus_count = hy_param.Choice('corpus_count', [5000, 10000,
15000, 20000])
    maxlen = hy_param.Choice('maxlen', [150, 200, 250, 300])
    embedding_features = hy_param.Choice('embedding_features', [150,
200, 250, 300])
    # learning rate param
    lr = hy_param.Choice('learning_rate', [1e-3, 5e-4])
    # regularization param
    regL1 = hy_param.Float('regularization1', 0.0, 0.1, step=0.005)
    regL2 = hy_param.Float('regularization2', 0.0, 0.1, step=0.005)
    # kernel initialization param
    kernel_initializers = hy_param.Choice('kernel_initializer',
['he_uniform', 'he_normal'])
    # dropout 1 param
    dropout1 = hy_param.Float('dropout1', 0, 0.5, step=0.2)
    # dropout 2 param
    dropout2 = hy_param.Float('dropout2', 0, 0.5, step=0.2)

    # Embedding
    x = Embedding(corpus_count, embedding_features, input_length=
maxlen)(inputs)

    # RNN
    RNN_layers = hy_param.Int('rnn_layers', 1, 3)
    for i in range(RNN_layers):
        if i < RNN_layers-1:
            seq = True
        else:
            seq = False
        i = str(i)
        filters = hy_param.Int('filters_' + i, 64, 260, step=64)
        rnn_dropout = hy_param.Float('dropout_rnn' + i, 0, 0.5,
step=0.2)
        rnn_type = hy_param.Choice('rnn_type_' + i, values=['lstm',
'gru'])
        if rnn_type == 'lstm':
            x = LSTM(filters, activation= 'relu', dropout=
rnn_dropout, return_sequences= seq)(x)
        elif rnn_type == 'gru':
            x = GRU(filters, activation= 'relu', dropout= rnn_dropout,
return_sequences= seq)(x)

    # Dropout 1
    x = Dropout(dropout1)(x)

```

```

# Batch normalization 1
x = BatchNormalization()(x)

# Dense 1 number of layers param
dense_layers1 = hy_param.Int('dense_layers1', 1, 3)

# Dense 1
for i in range(dense_layers1):
    dense_filters1 = 128 / (2 ** i)
    x = Dense(dense_filters1, activation='relu',
kernel_regularizer= L1L2(l1=regL1, l2=regL2),
kernel_initializer=kernel_initializers)(x)

# Dropout 2
x = Dropout(dropout2)(x)

# Batch normalization 2
x = BatchNormalization()(x)

# Dense 2 filter param
dense_filters2 = hy_param.Int('dense_layers2', 10, 21, step= 5)

# Dense 2
x = Dense(dense_filters2, activation='relu', kernel_regularizer=
L1L2(l1=regL1, l2=regL2), kernel_initializer=kernel_initializers)(x)

# Dense+softmax
outputs = Dense(cat_classes, activation='softmax',
kernel_regularizer= L1L2(l1=regL1, l2=regL2))(x)

# Model
model = Model(inputs, outputs)
model.compile(loss='categorical_crossentropy',
optimizer= Adam(lr),
metrics=[F1Score(num_classes= 3, average= 'macro')])

return model

tuner = RandomSearch(modelBuilder, objective='val_loss',
max_trials=trials,
project_name='ECommerce RNN Result')

tuner.search_space_summary()

Search space summary
Default search space size: 15
corpus_count (Choice)
{'default': 5000, 'conditions': [], 'values': [5000, 10000, 15000,
20000], 'ordered': True}
maxlen (Choice)
{'default': 150, 'conditions': [], 'values': [150, 200, 250, 300],

```

```

'ordered': True}
embedding_features (Choice)
{'default': 150, 'conditions': [], 'values': [150, 200, 250, 300],
'ordered': True}
learning_rate (Choice)
{'default': 0.001, 'conditions': [], 'values': [0.001, 0.0005],
'ordered': True}
regularization1 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.1,
'step': 0.005, 'sampling': None}
regularization2 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.1,
'step': 0.005, 'sampling': None}
kernel_initializer (Choice)
{'default': 'he_uniform', 'conditions': [], 'values': ['he_uniform',
'he_normal'], 'ordered': False}
dropout1 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5,
'step': 0.2, 'sampling': None}
dropout2 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5,
'step': 0.2, 'sampling': None}
rnn_layers (Int)
{'default': None, 'conditions': [], 'min_value': 1, 'max_value': 3,
'step': 1, 'sampling': None}
filters_0 (Int)
{'default': None, 'conditions': [], 'min_value': 64, 'max_value': 260,
'step': 64, 'sampling': None}
dropout_rnn0 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5,
'step': 0.2, 'sampling': None}
rnn_type_0 (Choice)
{'default': 'lstm', 'conditions': [], 'values': ['lstm', 'gru'],
'ordered': False}
dense_layers1 (Int)
{'default': None, 'conditions': [], 'min_value': 1, 'max_value': 3,
'step': 1, 'sampling': None}
dense_layers2 (Int)
{'default': None, 'conditions': [], 'min_value': 10, 'max_value': 21,
'step': 5, 'sampling': None}

callback = EarlyStopping(monitor='val_loss', patience=4)

tuner.search(x_train_seq, y_train_en, batch_size=batch_size,
epochs=epoch,
            validation_data=(x_test_seq, y_test_en),
callbacks=[callback])

Trial 3 Complete [00h 06m 04s]
val_loss: 9.555488586425781

```

Best val\_loss So Far: 1.0153387784957886  
Total elapsed time: 00h 18m 58s

tuner.results\_summary()

Results summary

Results in ./ECommerce RNN Result

Showing 10 best trials

<keras\_tuner.engine.objective.Objective object at 0x7ff45a768890>

Trial summary

Hyperparameters:

corpus\_count: 5000

maxlen: 300

embedding\_features: 250

learning\_rate: 0.0005

regularization1: 0.025

regularization2: 0.0

kernel\_initializer: he\_uniform

dropout1: 0.0

dropout2: 0.2

rnn\_layers: 1

filters\_0: 128

dropout\_rnn0: 0.4

rnn\_type\_0: gru

dense\_layers1: 1

dense\_layers2: 15

Score: 1.0153387784957886

Trial summary

Hyperparameters:

corpus\_count: 15000

maxlen: 150

embedding\_features: 300

learning\_rate: 0.0005

regularization1: 0.075

regularization2: 0.05

kernel\_initializer: he\_normal

dropout1: 0.0

dropout2: 0.4

rnn\_layers: 2

filters\_0: 128

dropout\_rnn0: 0.2

rnn\_type\_0: lstm

dense\_layers1: 1

dense\_layers2: 10

filters\_1: 64

dropout\_rnn1: 0.0

rnn\_type\_1: lstm

Score: 7.9572625160217285

Trial summary

Hyperparameters:

corpus\_count: 15000

```

maxlen: 150
embedding_features: 200
learning_rate: 0.001
regularization1: 0.03
regularization2: 0.085
kernel_initializer: he_uniform
dropout1: 0.0
dropout2: 0.2
rnn_layers: 3
filters_0: 64
dropout_rnn0: 0.0
rnn_type_0: gru
dense_layers1: 2
dense_layers2: 15
filters_1: 128
dropout_rnn1: 0.0
rnn_type_1: lstm
filters_2: 64
dropout_rnn2: 0.0
rnn_type_2: lstm
Score: 9.555488586425781

```

```

best_neural_model = tuner.get_best_models(num_models=1)[0]
best_neural_model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 200)]	0
embedding (Embedding)	(None, 200, 250)	1250000
gru (GRU)	(None, 128)	145920
dropout (Dropout)	(None, 128)	0
batch_normalization (Batch Normalization)	(None, 128)	512
dense (Dense)	(None, 128)	16512
dropout_1 (Dropout)	(None, 128)	0
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dense_1 (Dense)	(None, 15)	1935
dense_2 (Dense)	(None, 3)	48

```
=====
Total params: 1,415,439
Trainable params: 1,414,927
Non-trainable params: 512
=====
```

```
y_pred_train = getSentiment(best_neural_model.predict(x_train_seq))
y_pred_test = getSentiment(best_neural_model.predict(x_test_seq))
```

```
100/100 [=====] - 10s 102ms/step
25/25 [=====] - 3s 102ms/step
```

```
getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,
classes, 'Best Neural Net')
```

```
accuracy score train 0.9371875
accuracy score test 0.9375
```

Train classification report: Best Neural Net

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	74
Neutral	0.00	0.00	0.00	127
Positive	0.94	1.00	0.97	2999
accuracy			0.94	3200
macro avg	0.31	0.33	0.32	3200
weighted avg	0.88	0.94	0.91	3200

Test classification report: Best Neural Net

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	19
Neutral	0.00	0.00	0.00	31
Positive	0.94	1.00	0.97	750
accuracy			0.94	800
macro avg	0.31	0.33	0.32	800
weighted avg	0.88	0.94	0.91	800

Train confusion matrix: Best Neural Net

	Negative	Neutral	Positive
Negative	0	0	74
Neutral	0	0	127
Positive	0	0	2999



Test confusion matrix: Best Neural Net


	Negative	Neutral	Positive
Negative	0	0	19
Neutral	0	0	31
Positive	0	0	750

roc auc score train class Negative: 0.5


roc auc score train class Neutral: 0.5

roc auc score train class Positive: 0.5


Best Neural Net Train Receiver Operating Characteristic: Negative class



Best Neural Net Train Receiver Operating Characteristic: Neutral class



Best Neural Net Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.5


roc auc score test class Neutral: 0.5

roc auc score test class Positive: 0.5


Best Neural Net Test Receiver Operating Characteristic: Negative class



Best Neural Net Test Receiver Operating Characteristic: Neutral class



Best Neural Net Test Receiver Operating Characteristic: Positive class



```
test_y_pred = getSentiment(best_neural_model.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 4s 103ms/step
Test accuracy score 0.937
```

Classification report

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	24
Neutral	0.00	0.00	0.00	39
Positive	0.94	1.00	0.97	937
accuracy			0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```
best_neural_model.save('best_neural_net_random')
```

```
tuner = BayesianOptimization(modelBuilder, objective='val_loss',
max_trials=trials,
project_name='ECommerce RNN Result Bayesian')
```

```
tuner.search_space_summary()
```

## Search space summary

Default search space size: 15

corpus\_count (Choice)

{'default': 5000, 'conditions': [], 'values': [5000, 10000, 15000, 20000], 'ordered': True}

maxlen (Choice)

{'default': 150, 'conditions': [], 'values': [150, 200, 250, 300], 'ordered': True}

embedding\_features (Choice)

{'default': 150, 'conditions': [], 'values': [150, 200, 250, 300], 'ordered': True}

learning\_rate (Choice)

{'default': 0.001, 'conditions': [], 'values': [0.001, 0.0005], 'ordered': True}

regularization1 (Float)

{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.1, 'step': 0.005, 'sampling': None}

regularization2 (Float)

{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.1, 'step': 0.005, 'sampling': None}

kernel\_initializer (Choice)

{'default': 'he\_uniform', 'conditions': [], 'values': ['he\_uniform', 'he\_normal'], 'ordered': False}

dropout1 (Float)

{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.5, 'step': 0.2, 'sampling': None}

dropout2 (Float)

{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.5, 'step': 0.2, 'sampling': None}

rnn\_layers (Int)

{'default': None, 'conditions': [], 'min\_value': 1, 'max\_value': 3, 'step': 1, 'sampling': None}

filters\_0 (Int)

{'default': None, 'conditions': [], 'min\_value': 64, 'max\_value': 260, 'step': 64, 'sampling': None}

dropout\_rnn0 (Float)

{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.5, 'step': 0.2, 'sampling': None}

rnn\_type\_0 (Choice)

{'default': 'lstm', 'conditions': [], 'values': ['lstm', 'gru'], 'ordered': False}

dense\_layers1 (Int)

{'default': None, 'conditions': [], 'min\_value': 1, 'max\_value': 3, 'step': 1, 'sampling': None}

dense\_layers2 (Int)

{'default': None, 'conditions': [], 'min\_value': 10, 'max\_value': 21, 'step': 5, 'sampling': None}

tuner.search(x\_train\_seq, y\_train\_en, batch\_size=batch\_size,  
epochs=epoch,

```
        validation_data=(x_test_seq, y_test_en),  
callbacks=[callback])
```

```
Trial 3 Complete [00h 09m 17s]  
val_loss: nan
```

```
Best val_loss So Far: 0.6675423383712769  
Total elapsed time: 00h 21m 27s
```

```
tuner.results_summary()
```

```
Results summary
```

```
Results in ./ECommerce RNN Result Bayesian
```

```
Showing 10 best trials
```

```
<keras_tuner.engine.objective.Objective object at 0x7ff463a1c450>
```

```
Trial summary
```

```
Hyperparameters:
```

```
corpus_count: 15000
```

```
maxlen: 300
```

```
embedding_features: 300
```

```
learning_rate: 0.001
```

```
regularization1: 0.05
```

```
regularization2: 0.03
```

```
kernel_initializer: he_normal
```

```
dropout1: 0.2
```

```
dropout2: 0.0
```

```
rnn_layers: 1
```

```
filters_0: 128
```

```
dropout_rnn0: 0.4
```

```
rnn_type_0: lstm
```

```
dense_layers1: 1
```

```
dense_layers2: 20
```

```
filters_1: 192
```

```
dropout_rnn1: 0.4
```

```
rnn_type_1: gru
```

```
filters_2: 192
```

```
dropout_rnn2: 0.4
```

```
rnn_type_2: gru
```

```
Score: 0.6675423383712769
```

```
Trial summary
```

```
Hyperparameters:
```

```
corpus_count: 20000
```

```
maxlen: 300
```

```
embedding_features: 250
```

```
learning_rate: 0.0005
```

```
regularization1: 0.09
```

```
regularization2: 0.1
```

```
kernel_initializer: he_uniform
```

```
dropout1: 0.2
```

```
dropout2: 0.4
```

```
rnn_layers: 3
```

```

filters_0: 64
dropout_rnn0: 0.4
rnn_type_0: gru
dense_layers1: 3
dense_layers2: 15
filters_1: 64
dropout_rnn1: 0.0
rnn_type_1: lstm
filters_2: 64
dropout_rnn2: 0.0
rnn_type_2: lstm
Score: 9.053215026855469
Trial summary
Hyperparameters:
corpus_count: 5000
maxlen: 300
embedding_features: 300
learning_rate: 0.001
regularization1: 0.0
regularization2: 0.0
kernel_initializer: he_normal
dropout1: 0.2
dropout2: 0.0
rnn_layers: 1
filters_0: 256
dropout_rnn0: 0.4
rnn_type_0: lstm
dense_layers1: 1
dense_layers2: 20
filters_1: 256
dropout_rnn1: 0.4
rnn_type_1: gru
filters_2: 256
dropout_rnn2: 0.4
rnn_type_2: gru
Score: nan

```

```

best_neural_model = tuner.get_best_models(num_models=1)[0]
best_neural_model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 200)]	0
embedding (Embedding)	(None, 200, 300)	4500000
lstm (LSTM)	(None, 128)	219648
dropout (Dropout)	(None, 128)	0

batch_normalization (Batch Normalization)	(None, 128)	512
dense (Dense)	(None, 128)	16512
dropout_1 (Dropout)	(None, 128)	0
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dense_1 (Dense)	(None, 20)	2580
dense_2 (Dense)	(None, 3)	63

```

=====
Total params: 4,739,827
Trainable params: 4,739,315
Non-trainable params: 512
=====

```

```

y_pred_train = getSentiment(best_neural_model.predict(x_train_seq))
y_pred_test = getSentiment(best_neural_model.predict(x_test_seq))

100/100 [=====] - 15s 147ms/step
25/25 [=====] - 4s 144ms/step

getPerformance(y_train_label, y_pred_train, y_test_label, y_pred_test,
classes, 'Best Neural Net (Bayesian)')

accuracy score train 0.9371875
accuracy score test 0.9375

```

Train classification report: Best Neural Net (Bayesian)

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	74
Neutral	0.00	0.00	0.00	127
Positive	0.94	1.00	0.97	2999
accuracy			0.94	3200
macro avg	0.31	0.33	0.32	3200
weighted avg	0.88	0.94	0.91	3200

Test classification report: Best Neural Net (Bayesian)

	precision	recall	f1-score	support
--	-----------	--------	----------	---------



Negative	0.00	0.00	0.00	19
Neutral	0.00	0.00	0.00	31
Positive	0.94	1.00	0.97	750
accuracy			0.94	800
macro avg	0.31	0.33	0.32	800
weighted avg	0.88	0.94	0.91	800

Train confusion matrix: Best Neural Net (Bayesian)

	Negative	Neutral	Positive
Negative	0	0	74
Neutral	0	0	127
Positive	0	0	2999

Test confusion matrix: Best Neural Net (Bayesian)


	Negative	Neutral	Positive
Negative	0	0	19
Neutral	0	0	31
Positive	0	0	750

roc auc score train class Negative: 0.5


roc auc score train class Neutral: 0.5

roc auc score train class Positive: 0.5


Best Neural Net (Bayesian) Train Receiver Operating Characteristic: Negative class



Best Neural Net (Bayesian) Train Receiver Operating Characteristic: Neutral class



Best Neural Net (Bayesian) Train Receiver Operating Characteristic: Positive class




roc auc score test class Negative: 0.5


roc auc score test class Neutral: 0.5

roc auc score test class Positive: 0.5


Best Neural Net (Bayesian) Test Receiver Operating Characteristic: Negative class



Best Neural Net (Bayesian) Test Receiver Operating Characteristic: Neutral class



Best Neural Net (Bayesian) Test Receiver Operating Characteristic: Positive class



```
test_y_pred = getSentiment(best_neural_model.predict(test_X_seq))
print('Test accuracy score {} \n'.format(accuracy_score(test_y,
test_y_pred)))
print('Classification report\n')
print(classification_report(test_y, test_y_pred))
```

```
32/32 [=====] - 5s 143ms/step
Test accuracy score 0.937
```

Classification report

	precision	recall	f1-score	support
Negative	0.00	0.00	0.00	24
Neutral	0.00	0.00	0.00	39
Positive	0.94	1.00	0.97	937
accuracy			0.94	1000
macro avg	0.31	0.33	0.32	1000
weighted avg	0.88	0.94	0.91	1000

```
best_neural_model.save('best_neural_net_bayesian')
```

Still, by looking that the perormance, LSTM (w2v\_lstm\_model) is performing the best.

## Topic Modeling

### 7. Cluster similar reviews.

```
ecomp_train_text = pd.read_pickle("ecom_train_processed.pkl")
ecom_train_text = ecomp_train_text[['reviews.clean_text',
'sentiment']]

ecom_train_text['sentiment'].value_counts()

Positive      3749
Neutral       158
Negative        93
Name: sentiment, dtype: int64

text_pos = ecomp_train_text[ecom_train_text['sentiment'] ==
'Positive'].drop(['sentiment'], axis=1)[:93]
text_neg = ecomp_train_text[ecom_train_text['sentiment'] ==
'Negative'].drop(['sentiment'], axis=1)[:93]
text_neu = ecomp_train_text[ecom_train_text['sentiment'] ==
'Neutral'].drop(['sentiment'], axis=1)[:93]

X_text = pd.concat([text_pos, text_neg, text_neu], axis=0)
X_text.shape


(279, 1)

# get tfidf feature vectors
tfidf = TfidfVectorizer(max_features = 5000, ngram_range=(1, 2))
X_text_vec = tfidf.fit_transform(X_text.iloc[:, 0])

We can use KMeans clustering here.

cost = []
for i in range(1, 10):
    KM = KMeans(n_clusters = i, max_iter = 500)
    KM.fit(X_text_vec)
    cost.append(KM.inertia_)

plt.plot(range(1, 10), cost, color='g', linewidth='3')
plt.xlabel("Value of K")
plt.ylabel("Squared Error (Cost)")
plt.show()
```



We can choose 5 as optimum K value for Kmeans, because we are not seeing any elbow point.

```
KM = KMeans(n_clusters = 5, max_iter = 500)
labels = KM.fit_predict(X_text_vec)
```

We can view the result with the help of PCA.


```
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_text_vec.toarray())
df = pd.DataFrame(X_pca, columns=['PC1', 'PC2'])
df['Text'] = X_text.iloc[:, 0]
df['label'] = labels
df.shape
(279, 4)
df[:1]
```

	PC1	PC2	
Text \			
0	-0.019831	0.012978	purchas black fridaypro great price even salev...
label			
0	4		

```

fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('PC 1', fontsize = 15)
ax.set_ylabel('PC 2', fontsize = 15)
ax.set_title('Review PCA with cluster label', fontsize = 20)
targets = sorted(list(df['label'].unique()))
for target in targets:
    idx = df['label'] == target
    ax.plot(df.loc[idx, 'PC1'], df.loc[idx, 'PC2'], linestyle='none',
marker='o', label=target)
ax.legend(targets)
ax.grid()

```



Lets check review texts for different clusters

```

for i in range(5):
    print('cluster ' + str(i))
    text_i = df[df['label'] == i]['Text']
    print(text_i[:1].values)
    print(text_i[1:2].values)
    print('-----')

cluster 0
['purchas two amazon echo plus two dot plus four fire stick hub philip
hue lamp famili christma happi purchas learn much alexa start daili
routin alexa program whatev would like includ news weather music
horoscop also start day compliment think import alexa gave best chili
recip mean best call chili want husband use alexa stay organ busi date
remind way go']
['replac fire kindl love hd qualiti']
-----
cluster 1
['good product exact want good price']
['great product light weight wish wifi download onlin']
-----
cluster 2
['averag alexa option show thing screen still limit']
['rd one ive purchas ive bought one niec case compar one held protect
tablet mani time drop']
-----
cluster 3
['cheap run chrome stuff return store']
['prime almost must keep amazon stuff one place']
-----
cluster 4
['purchas black fridaypro great price even saleveri power fast quad
core processor amaz soundwel builtcon amazon ad amazon need subsid
tablet remov add pay inabl access app except one amazon way abl
accomplish add googl play storenet great tablet money']
['got model sale price built hub control hue strip light andi want say
googl time sound like babi talk melol']
-----

```

We can use the preprocessed texts and do topic modelling on top of that using gensim library.

First we need to build a text corpus and use LDA model for topic modelling.

```

ecom_train_text = pd.read_pickle("ecom_train_processed.pkl")
cleaned_texts = ecom_train_text['reviews.clean_text'].values

corpus = [text.split() for text in cleaned_texts]
print('total texts', len(corpus))
print(corpus[0][:10])

```



```

total texts 4000
['purchas', 'black', 'fridaypro', 'great', 'price', 'even',
'saleveri', 'power', 'fast', 'quad']

dictionary = corpora.Dictionary(corpus)
print(dictionary)

Dictionary(3697 unique tokens: ['abl', 'access', 'accomplish', 'ad',
'add']...)

doc_matrix = [dictionary.doc2bow(text) for text in corpus]

num_topic = 10
lda = LdaModel(doc_matrix, num_topics= num_topic, id2word= dictionary,
passes= 40)

topics = lda.show_topics()
for topic in topics:
    print(topic)
    print('-----')

(0, '0.063*"love" + 0.036*"bought" + 0.030*"old" + 0.029*"year" +
0.026*"kindl" + 0.024*"tablet" + 0.024*"use" + 0.022*"one" +
0.019*"easi" + 0.013*"gift"')
-----
(1, '0.029*"kindl" + 0.028*"one" + 0.024*"purchas" + 0.021*"read" +
0.017*"book" + 0.012*"love" + 0.012*"new" + 0.012*"small" +
0.012*"bought" + 0.011*"work"')
-----
(2, '0.031*"read" + 0.027*"devic" + 0.027*"book" + 0.023*"amazon" +
0.019*"great" + 0.019*"fire" + 0.016*"kindl" + 0.016*"best" +
0.013*"size" + 0.012*"buy"')
-----
(3, '0.030*"use" + 0.024*"screen" + 0.022*"easi" + 0.020*"sound" +
0.017*"good" + 0.015*"speaker" + 0.013*"like" + 0.013*"light" +
0.012*"qualiti" + 0.011*"get"')
-----
(4, '0.020*"amazon" + 0.020*"tablet" + 0.019*"would" + 0.017*"use" +
0.017*"buy" + 0.015*"product" + 0.015*"great" + 0.013*"recommend" +
0.013*"good" + 0.013*"set"')
-----
(5, '0.070*"great" + 0.057*"tablet" + 0.049*"price" + 0.029*"kid" +
0.022*"good" + 0.021*"need" + 0.019*"product" + 0.017*"recommend" +
0.016*"work" + 0.012*"like"')
-----
(6, '0.059*"tablet" + 0.045*"use" + 0.031*"app" + 0.028*"great" +
0.020*"game" + 0.019*"good" + 0.018*"play" + 0.018*"fire" +
0.014*"amazon" + 0.013*"work"')
-----
(7, '0.030*"love" + 0.029*"great" + 0.022*"echo" + 0.022*"tap" +
0.019*"gift" + 0.015*"learn" + 0.015*"buy" + 0.014*"best" +

```

```

0.014*"product" + 0.013*"light"')
-----
(8, '0.045*"batteri" + 0.031*"life" + 0.029*"read" + 0.020*"make" +
0.016*"much" + 0.015*"worth" + 0.013*"kindl" + 0.011*"last" +
0.010*"generat" + 0.010*"book"')
-----
(9, '0.052*"echo" + 0.034*"alexa" + 0.029*"show" + 0.026*"love" +
0.025*"use" + 0.025*"music" + 0.022*"great" + 0.019*"home" +
0.015*"devic" + 0.015*"video"')
-----

```

Now lets create a dataframe to view these topics in a better way. We can choose top 15 values for each topic.

```
lda.show_topic(0, topn = 15)
```

```

[('love', 0.062940404),
 ('bought', 0.036242247),
 ('old', 0.02960274),
 ('year', 0.029061344),
 ('kindl', 0.026133668),
 ('tablet', 0.024342084),
 ('use', 0.024142405),
 ('one', 0.021713383),
 ('easi', 0.018683134),
 ('gift', 0.013292712),
 ('daughter', 0.013235307),
 ('game', 0.012386512),
 ('son', 0.012087541),
 ('fire', 0.011921115),
 ('purchas', 0.011557242)]

```

```

topic_dict = {}
for i in range(num_topic):
    words = lda.show_topic(i, topn = 15)
    topic_dict['Topic-'+str(i)] = [word[0] for word in words]

```

```

topic_df = pd.DataFrame(topic_dict)
topic_df

```

	Topic-0	Topic-1	Topic-2	Topic-3	Topic-4	Topic-5	Topic-6
0	love	kindl	read	use	amazon	great	tablet
1	bought	one	devic	screen	tablet	tablet	use
2	old	purchas	book	easi	would	price	app
3	year	read	amazon	sound	use	kid	great
4	kindl	book	great	good	buy	good	game

5	tablet	love	fire	speaker	product	need	good
6	use	new	kindl	like	great	product	play
7	one	small	best	light	recommend	recommend	fire
8	easi	bought	size	qualiti	good	work	amazon
9	gift	work	buy	get	set	like	work
10	daughter	fit	movi	great	easi	get	read
11	game	happi	better	want	work	friend	easi
12	son	would	price	read	one	lot	love
13	fire	buy	featur	better	time	would	need
14	purchas	back	love	turn	kid	enjoy	kid

	Topic-7	Topic-8	Topic-9
0	love	batteri	echo
1	great	life	alexa
2	echo	read	show
3	tap	make	love
4	gift	much	use
5	learn	worth	music
6	buy	kindl	great
7	best	last	home
8	product	generat	devic
9	light	book	video
10	amazon	reader	smart
11	like	th	play
12	go	light	amazon
13	hous	excel	dot
14	christma	great	ask

Lets display all the topics for better insights.


```
lda_display = pyLDAvis.gensim_models.prepare(lda, doc_matrix,
dictionary, sort_topics = False)
pyLDAvis.enable_notebook()
pyLDAvis.display(lda_display)

<IPython.core.display.HTML object>


pyLDAvis.save_html(lda_display, 'lda_display.html')
```

Ecommerce sentiment analysis was performed with ML and DL models, clustering and topic modeling.

Intertopic Distance Map (via multidimensional scaling)



Top-30 Most Salient Terms<sup>(1)</sup>



1. saliency(term w) = frequency(w) \* [sum\_t p(t | w) \* log(p(t | w)/p(t)))] for topics t; see Chuang et. al (2012)  
2. relevance(term w | topic t) = λ \* p(w | t) + (1 - λ) \* p(w | t)/p(w); see Sievert & Shirley (2014)